

# Using `ngx_lua` in UPYUN 2

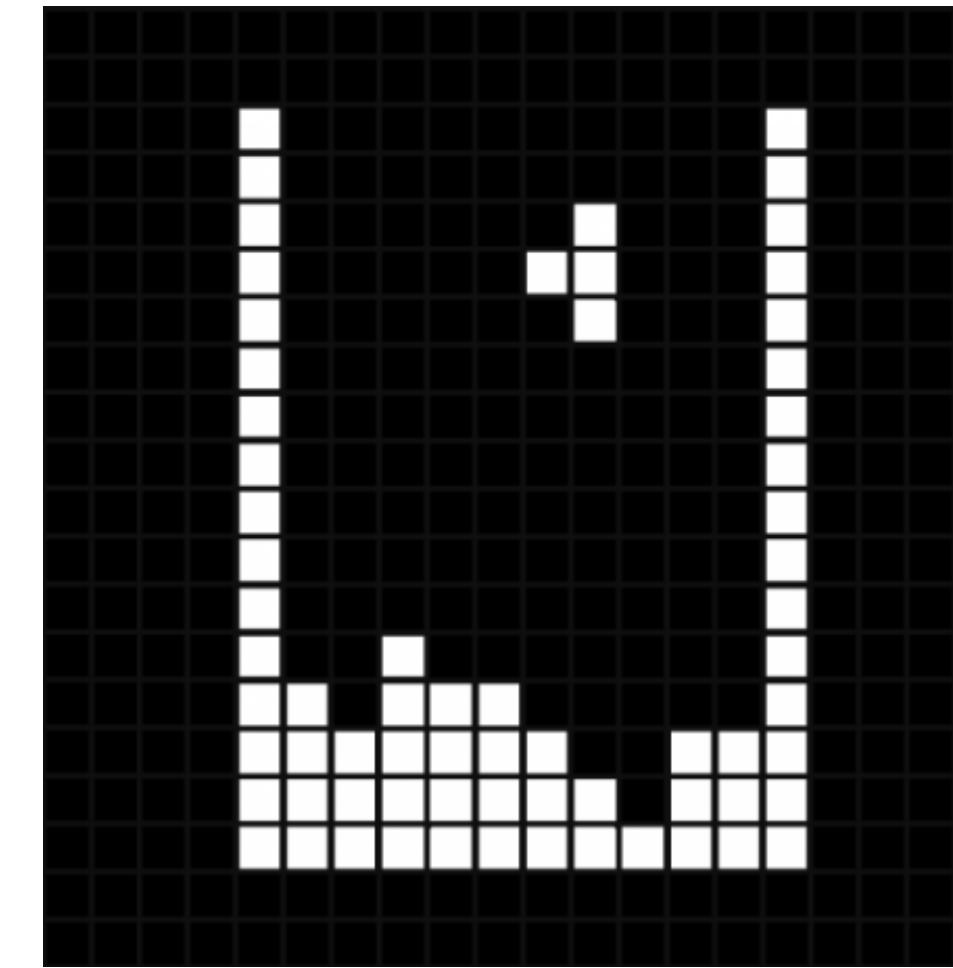


Monkey Zhang (timebug)  
2015.11 @ Beijing OpenResty Con





upyun



A Systems Engineer at **UPYUN**

★Email: [timebug.info@gmail.com](mailto:timebug.info@gmail.com)

★Github: <https://github.com/timebug>



```
$ ./configure --prefix=/opt/nginx \
--add-module=/path/to/lua-nginx-module
```

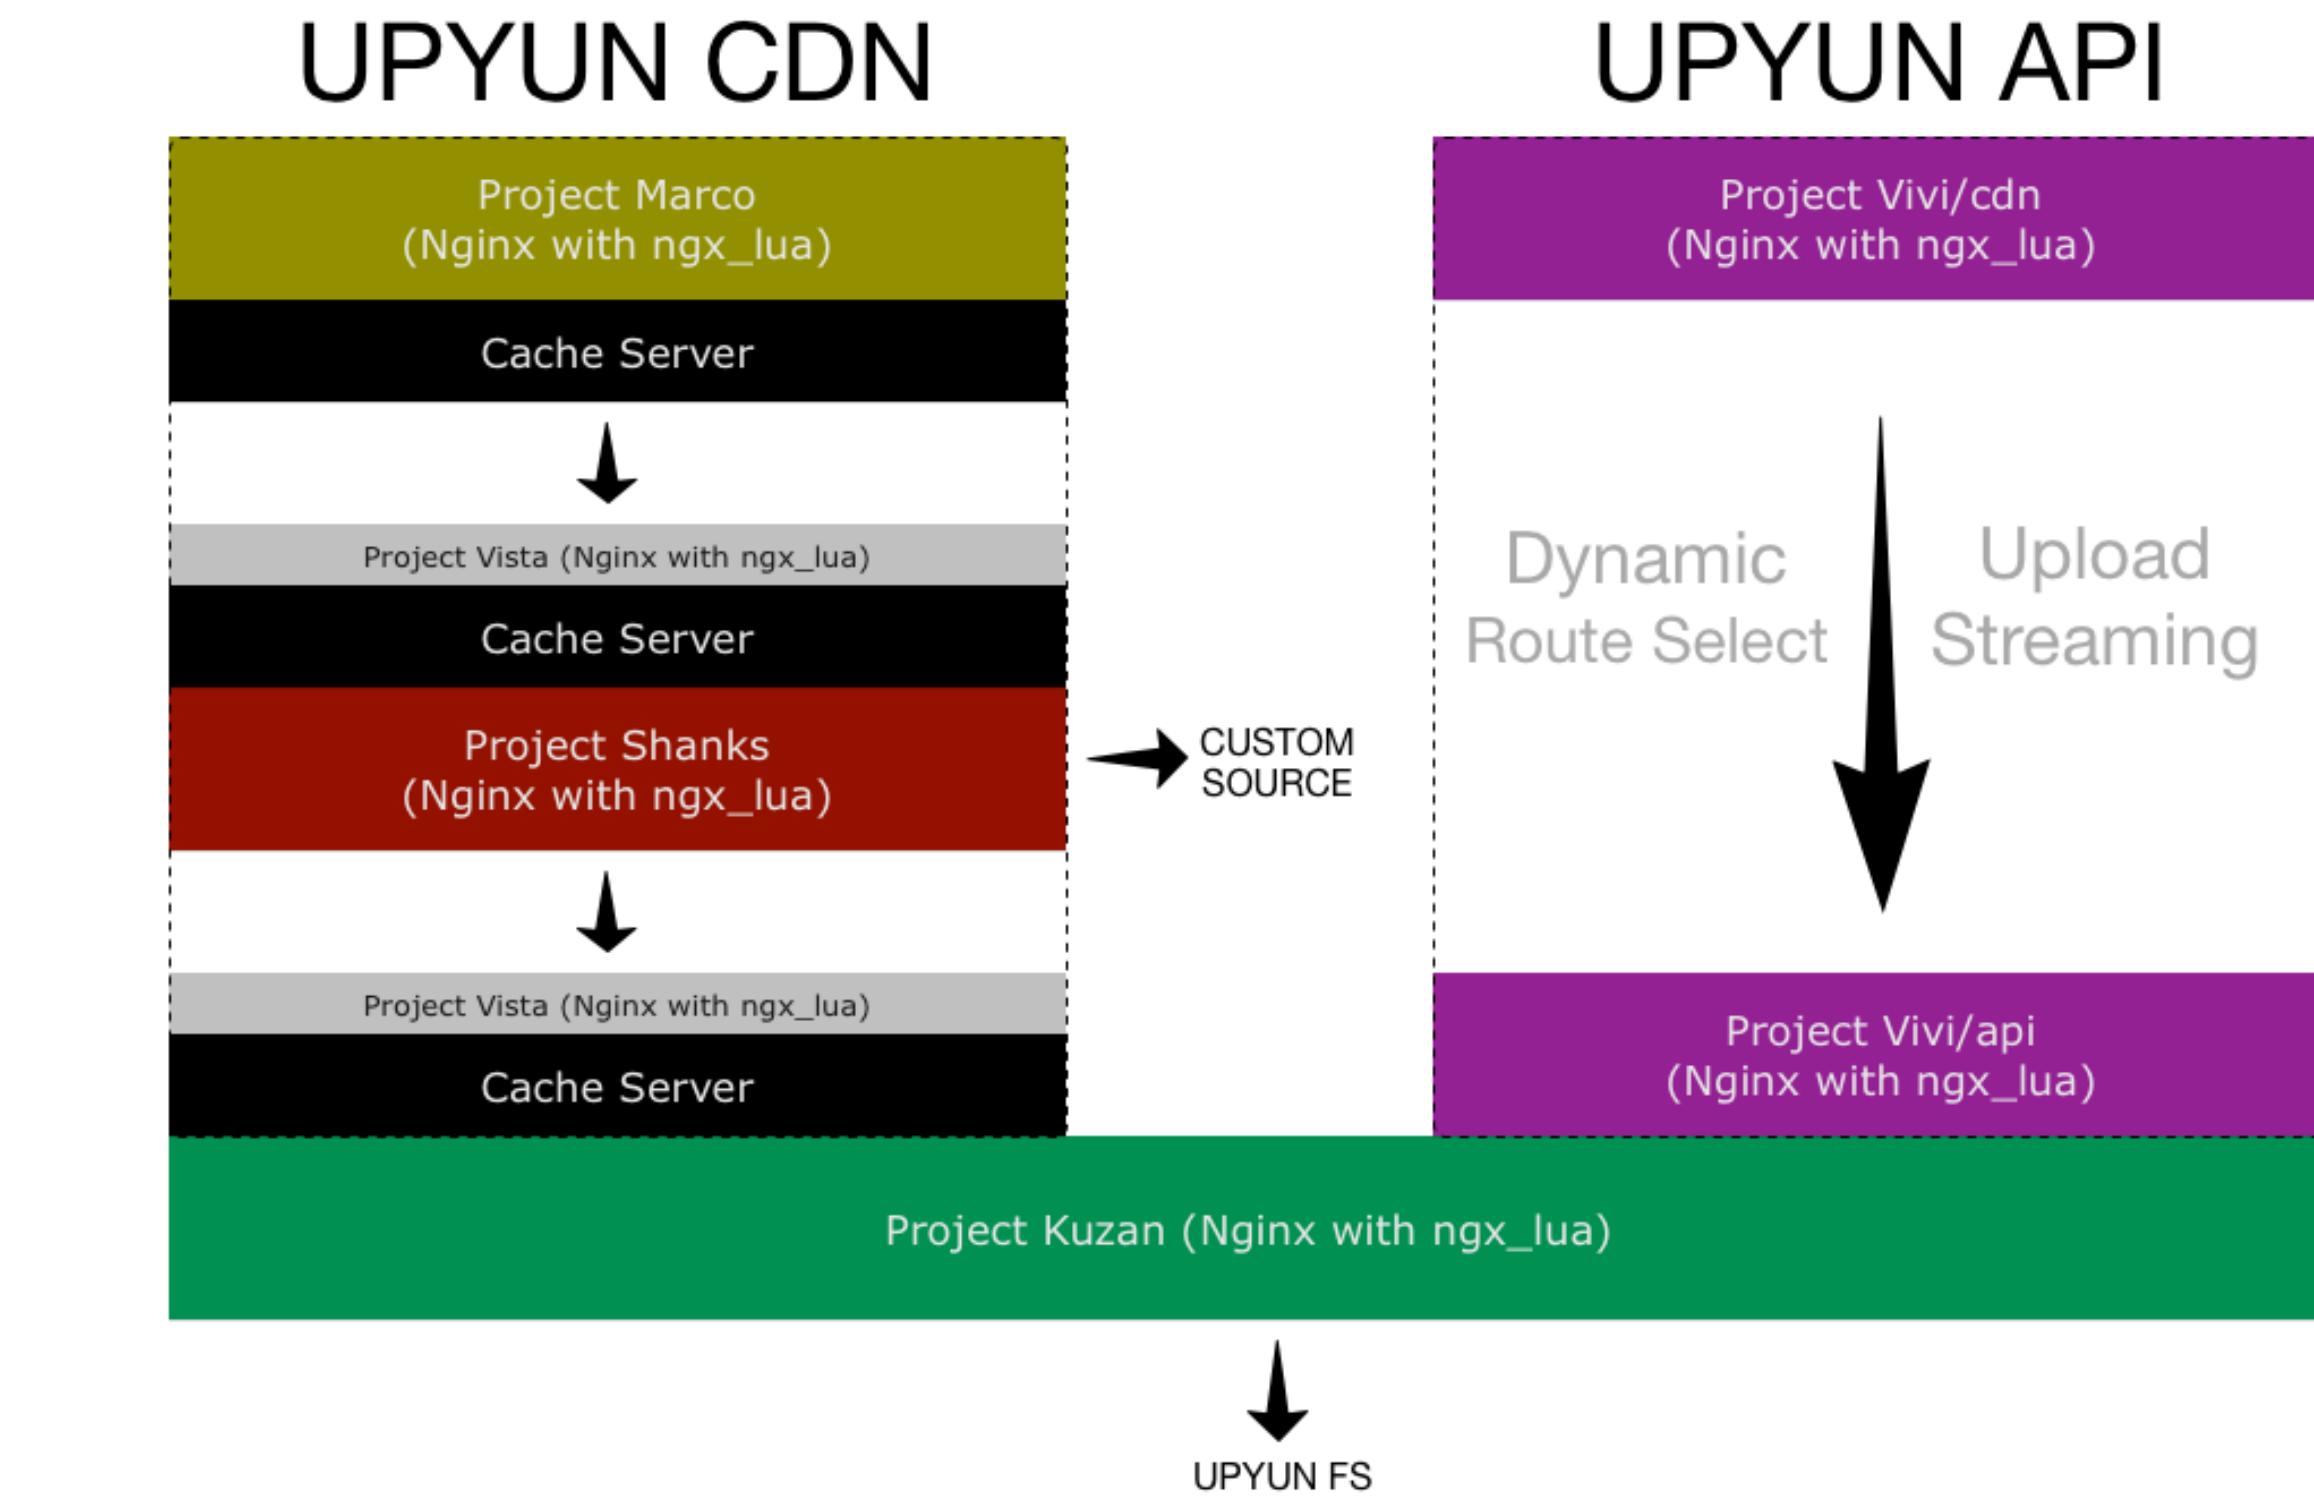
```
http {
    server {
        listen 8080;

        location /add {
            set $res '';

            rewrite_by_lua '
                local a = tonumber(ngx.var.arg_a) or 0
                local b = tonumber(ngx.var.arg_b) or 0
                ngx.var.res = a + b
            ';

            content_by_lua '
                ngx.say(ngx.var.res)
            ';
        }
    }
}
```

```
$ curl 'http://localhost:8080/add?a=6&b=7'
```



**UPYUN CDN & API** is built on top of  
**NGINX** with **ngx\_lua**

*Why not use OpenResty?*

# 40000+ lines **Lua**

lua-resty-sniff

[lua-resty-limit-req](#)

## lua-resty-httpipe

lua-resty-anticc

lua-resty-rewrite

[lua-resty-17monip](#)

**lua-resty-dbcache**

lua-resty-combo

lua-resty-httproxy

**lua-resty-checkups**

**lua-resty-argutils**

...

# Project Structure:

## NGINX with ngx\_lua

```
~/project/upyun/marco

├── Makefile
├── README.md
├── addons
│   └── ngx_upxxx_module
├── deps
├── nginx
│   ├── app
│   │   └── etc
│   │       └── config.lua
│   ├── lib
│   │   └── resty
│   │       └── httpipe.lua
│   └── src
│       ├── modules
│       │   └── referer.lua
│       ├── marco_init.lua
│       └── marco_log.lua
└── conf
    └── nginx.conf
patches
tests
util
└── deps
└── lua-releng
└── ver.cfg
```

.....  
**make install**

```
/usr/local/marco

├── luajit
└── nginx
    ├── app
    │   └── etc
    │       └── config.lua
    ├── lib
    │   └── resty
    │       └── httpipe.lua
    └── src
        ├── modules
        │   └── referer.lua
        ├── marco_init.lua
        └── marco_log.lua
    └── conf
        └── nginx.conf
    └── html
    └── logs
    └── sbin
        └── nginx
```

# Project Structure:

## Quick Start & Run

**make deps**

.....

util/ver.cfg

V\_PCRE=8.34  
V\_NGINX=1.7.10  
V\_LUAJIT=2.1-20150223  
V\_LUA\_CJSON=2.1.0  
V\_NGX\_DEVEL\_KIT=0.2.19  
V\_NGX\_LUA\_MODULE=0.9.15

**make configure**

.....

**make**

.....

**make install**

.....

### Makefile

```
INSTALL_LIBDIR=$(PREFIX)/nginx/app/lib/
configure: deps luajit
    @echo "==== Configuring Nginx $(V_NGINX) ===="
    cd $(NGINX_DIR) && ./configure \
        --with-pcre=$(ROOTDIR)/deps/pcre-$(V_PCRE) \
        --with-ld-opt="-Wl,-rpath,$(LUAJIT_LIB),-rpath,$(INSTALL_LIBDIR)" \
        --add-module=$(ROOTDIR)/deps/ngx-devel-kit-$(V_NGX_DEVEL_KIT) \
        --add-module=$(ROOTDIR)/deps/lua-nginx-module-$(V_NGX_LUA_MODULE) \
        --prefix=$(PREFIX)/nginx
    @echo "==== Successfully configure Nginx $(V_NGINX) ===="
```

# Project Structure:

## Development & Test

**make dev**

**make test**



Makefile

```
test:  
    util/lua-releng  
    py.test tests/test_marco.py
```

tests/test\_marco.py

```
class TestMarco(unittest.TestCase):  
  
    @no_error_log(["error"])  
    @grep_error_log(level=["info"],  
                    log_pattern="SSL_do_handshake[()[]] failed",  
                    log_out=["SSL_do_handshake() failed"])  
    def test_ssl_handler_no_certificate(self):  
        fake_resp = self.curl_ssl(sni="fake.com", verbose=True)  
        self.assertTrue("alert handshake failure" in fake_resp)
```

## nginx.conf

```
server_name *.b0.upaiyun.com
```

```
valid_referers, allow, deny
```

```
expires 7d
```

```
ssl_certificate* ssl_stapling*
```

```
upstream { server 127.0.0.1 }
```

```
max_fails=3 fail_timeout=30s  
health_check (*)
```

```
round-robin, ip_hash, hash (1.7.2+)
```

```
rewrite
```

```
...
```

## service

Custom Domain Binding

**Custom Antileech Rules and Redirect:**  
ip, user-agent, referer, token etc.

**Custom Cache Control:**  
support specific URI rules etc.

Custom SSL

Custom CDN Origin:  
support multi-network routing etc.

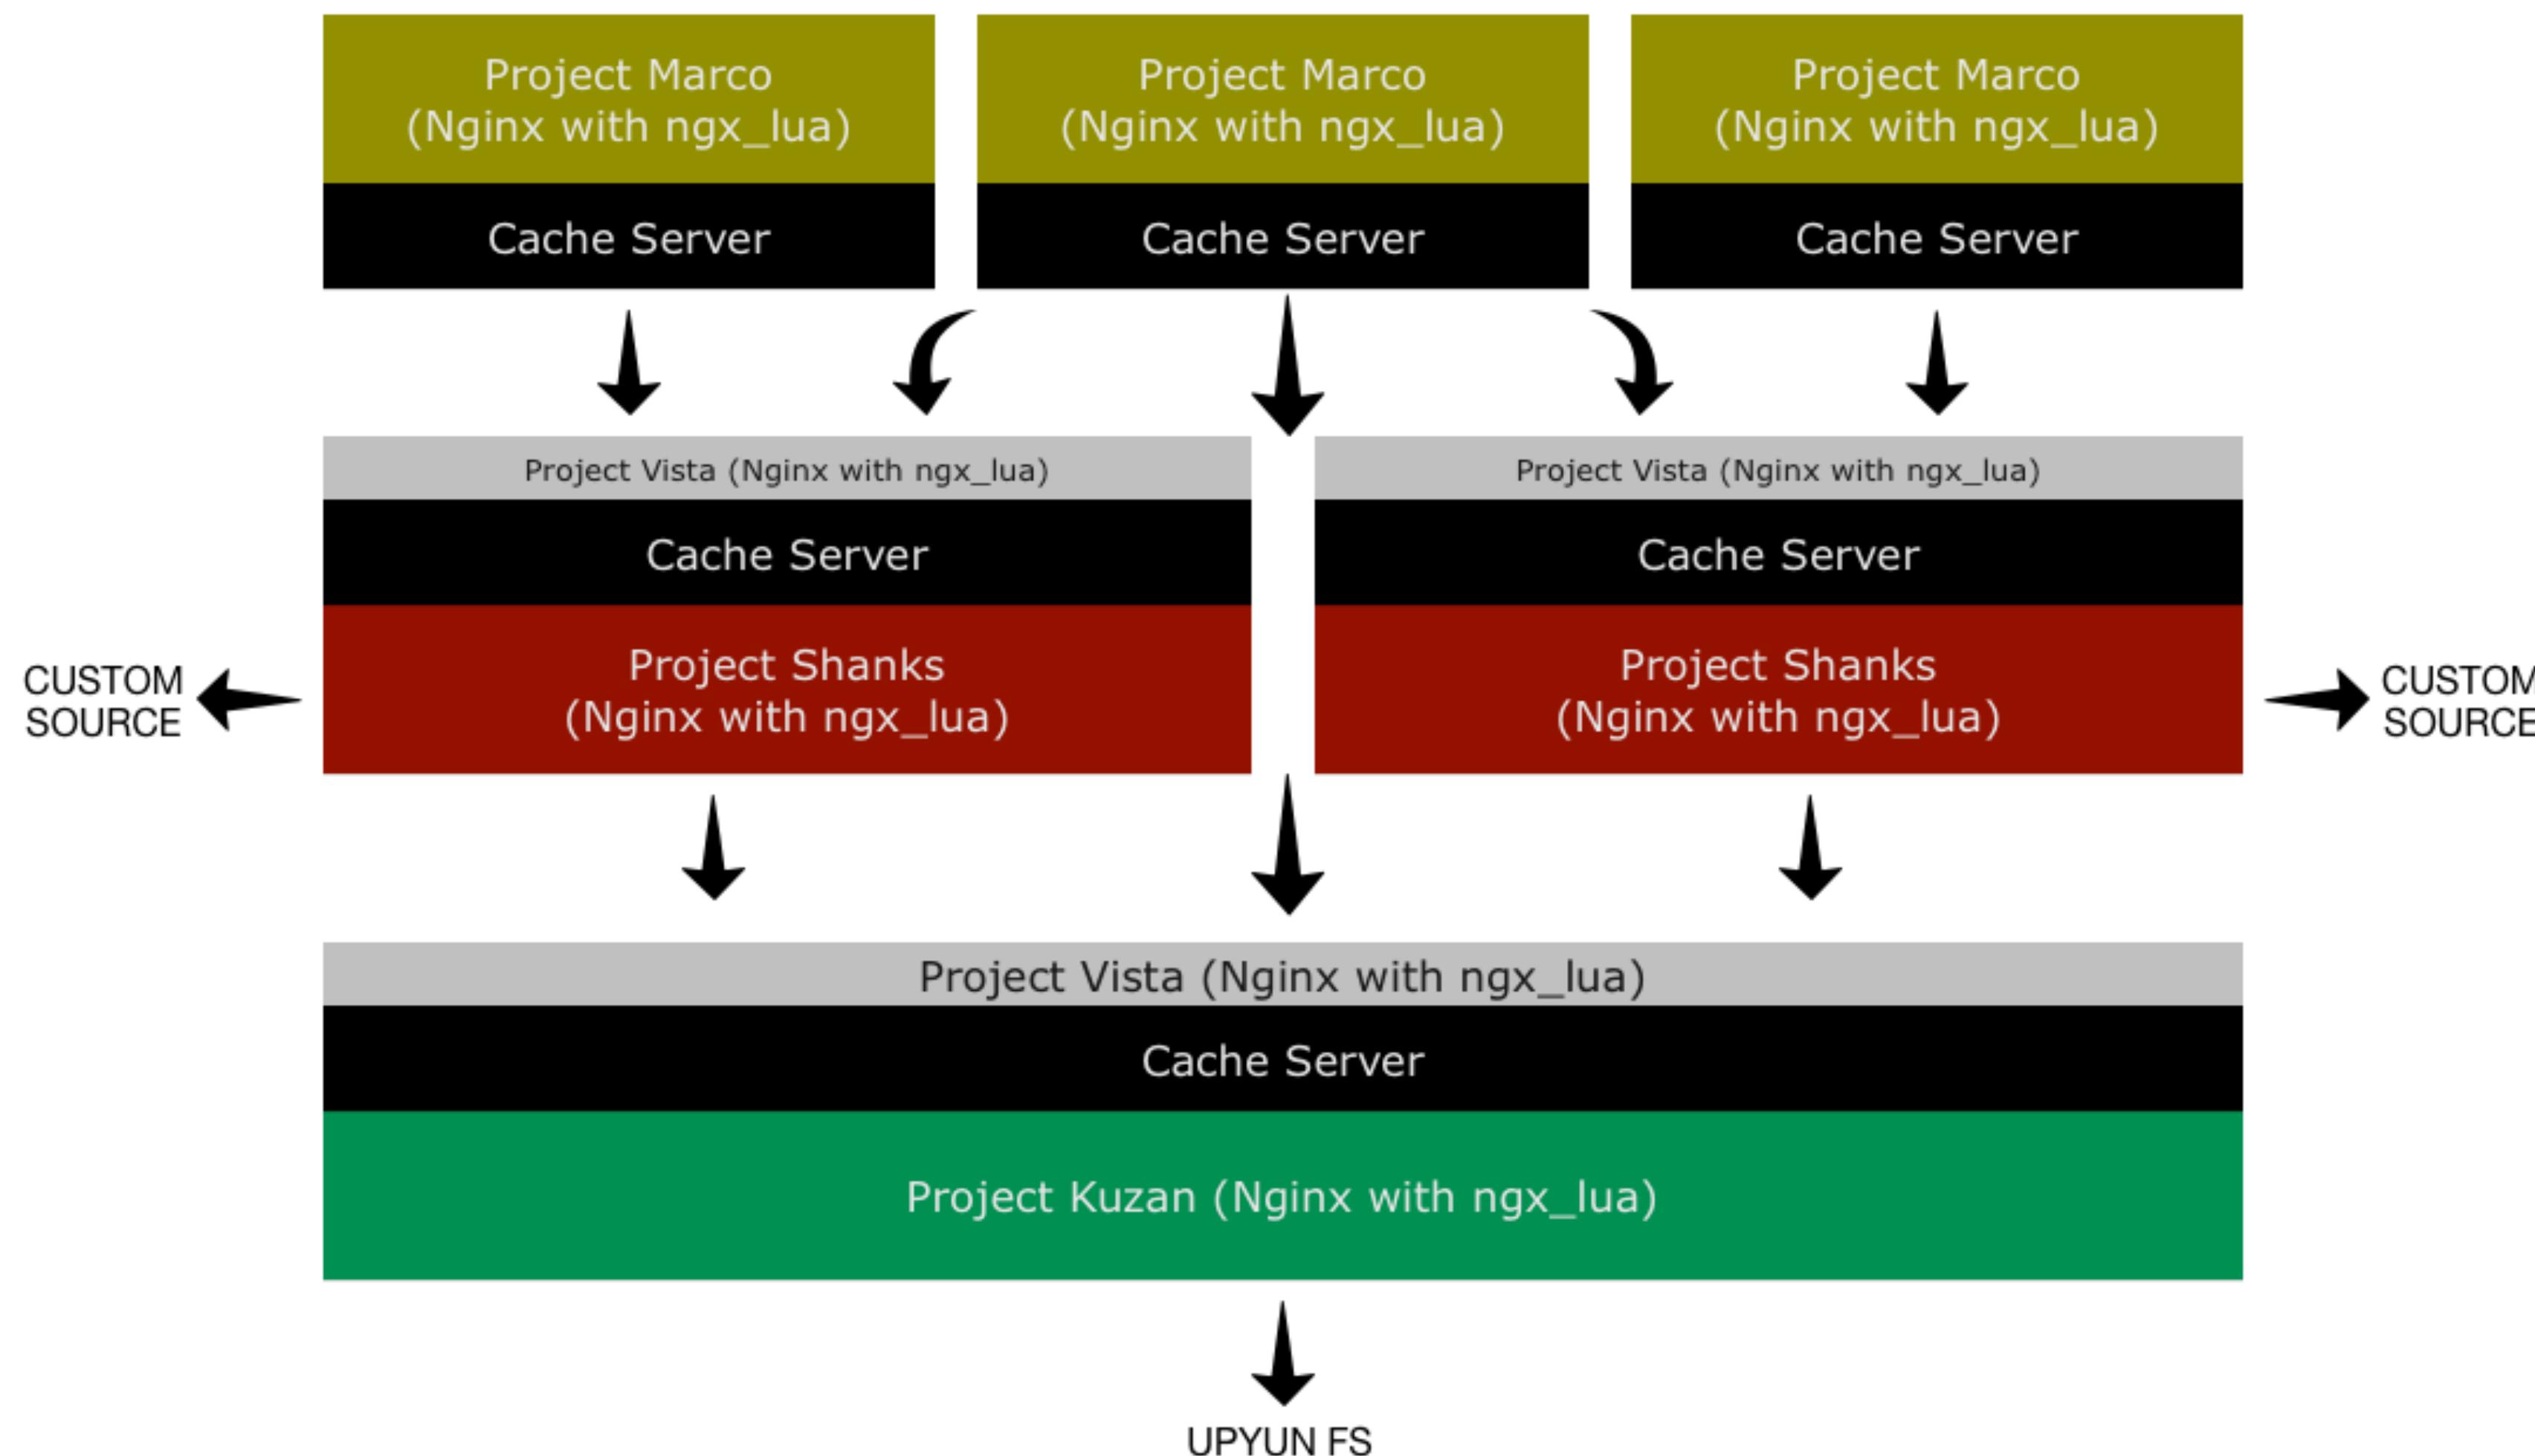
Custom Health Check Strategy:  
passive, active

Custom Load Balancing Strategy

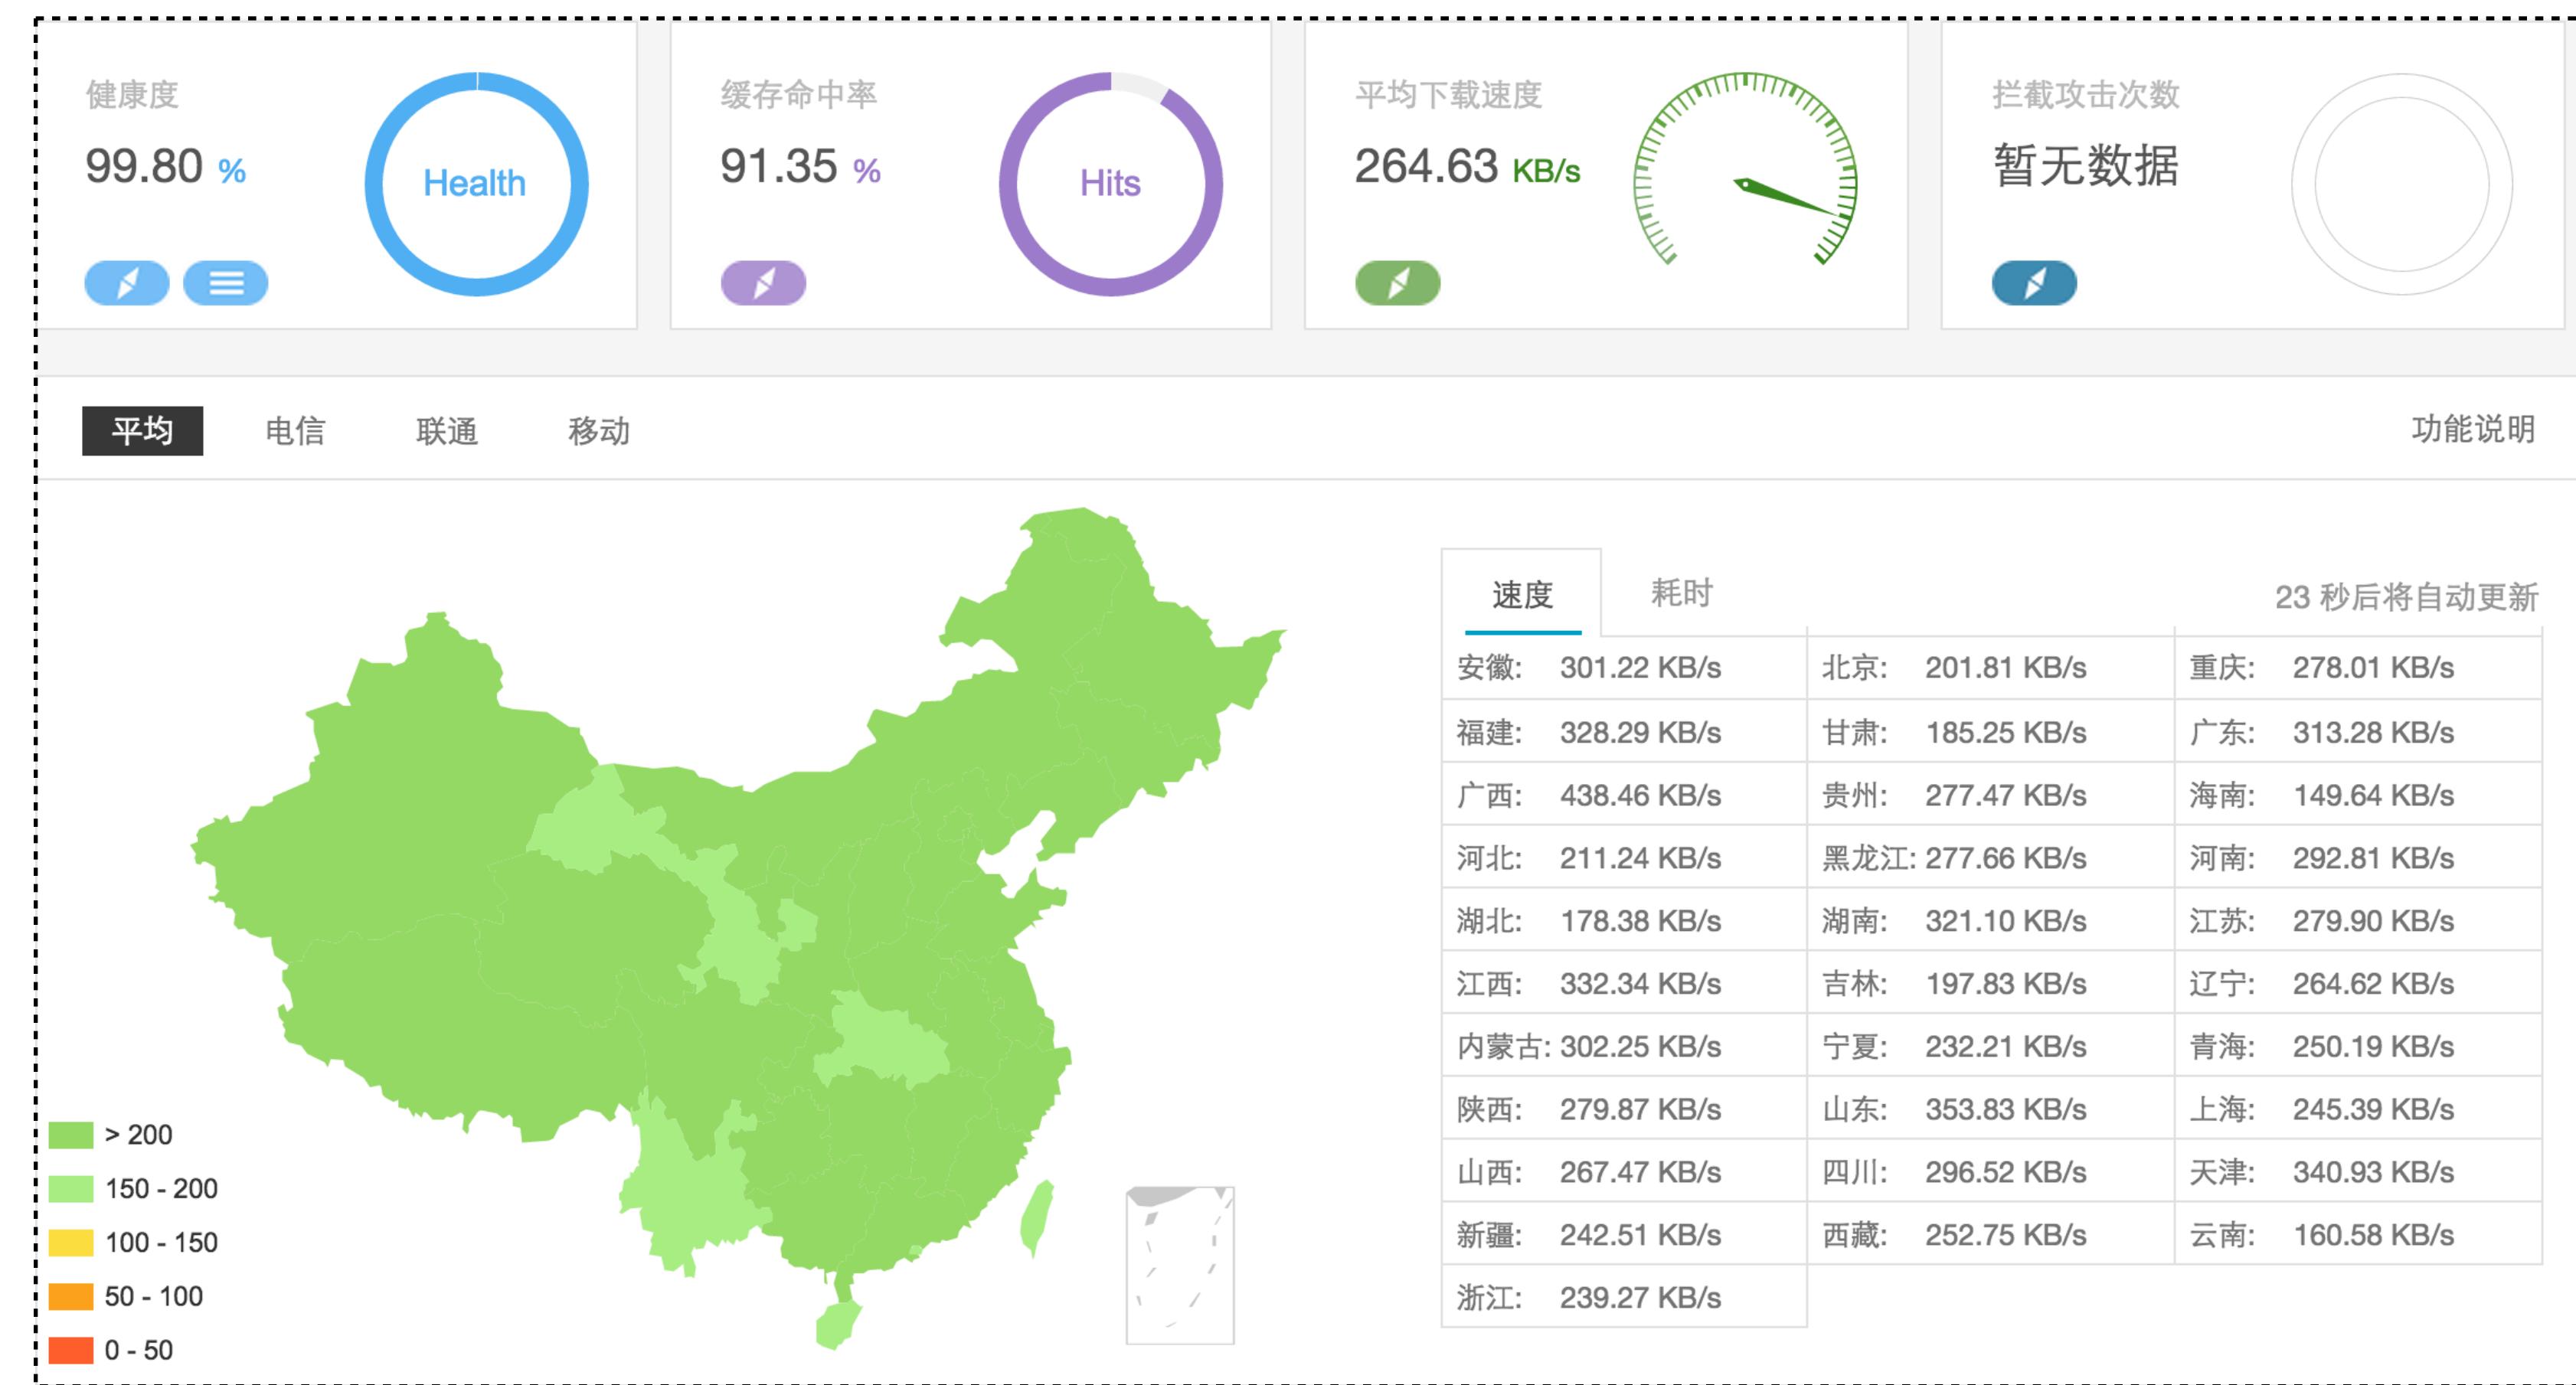
Custom URL rewrite

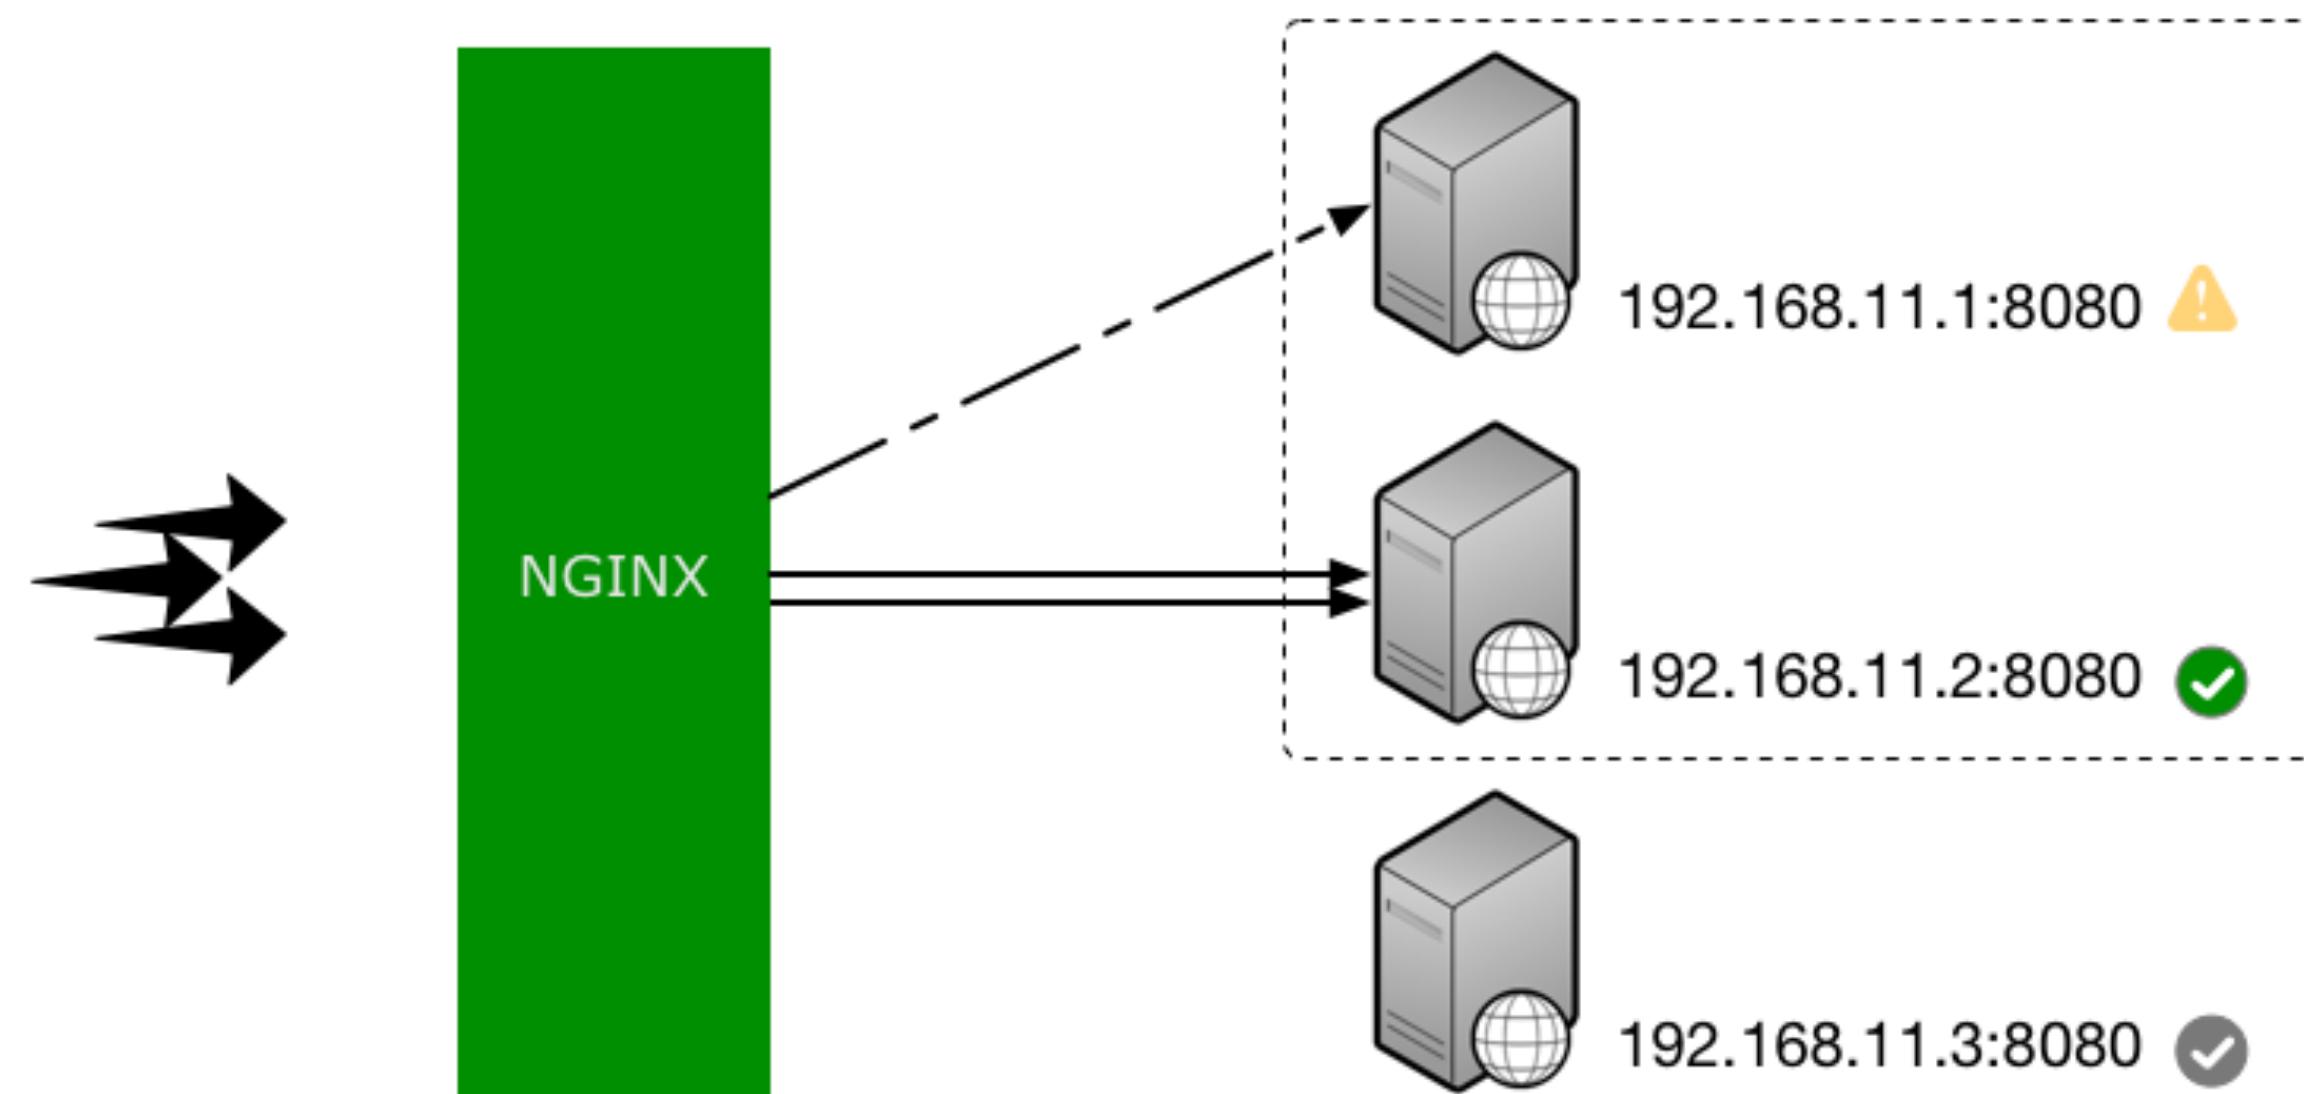
```
...
```

# UPYUN CDN



# 130+ Edge Nodes





```
upstream blog.upyun.com {  
    server 192.168.11.1:8080 weight=1 max_fails=10 fail_timeout=30s;  
    server 192.168.11.2:8080 weight=2 max_fails=10 fail_timeout=30s;  
  
    server 192.168.11.3:8080 weight=1 max_fails=10 fail_timeout=30s backup;  
  
    proxy_next_upstream error timeout http_500;  
    proxy_next_upstream_tries 2;  
}
```

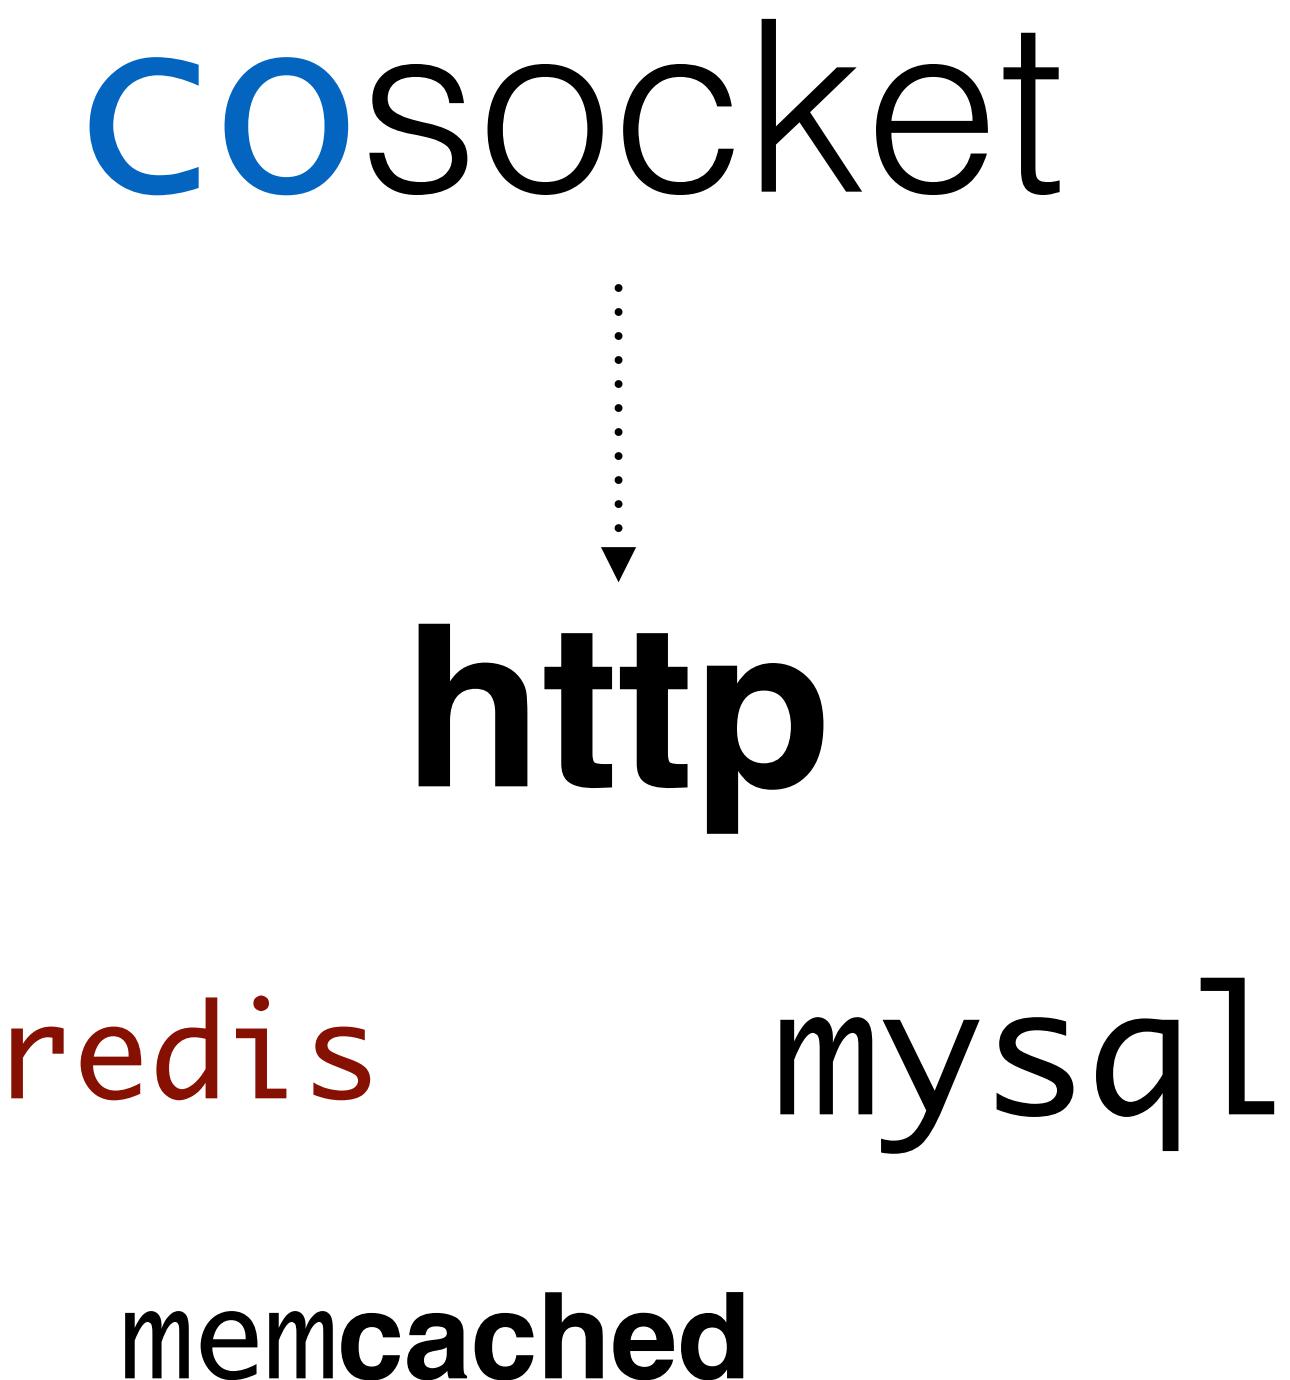
# Lua Upstream Configuration: lua-resty-checkups

```
-- app/etc/config.lua

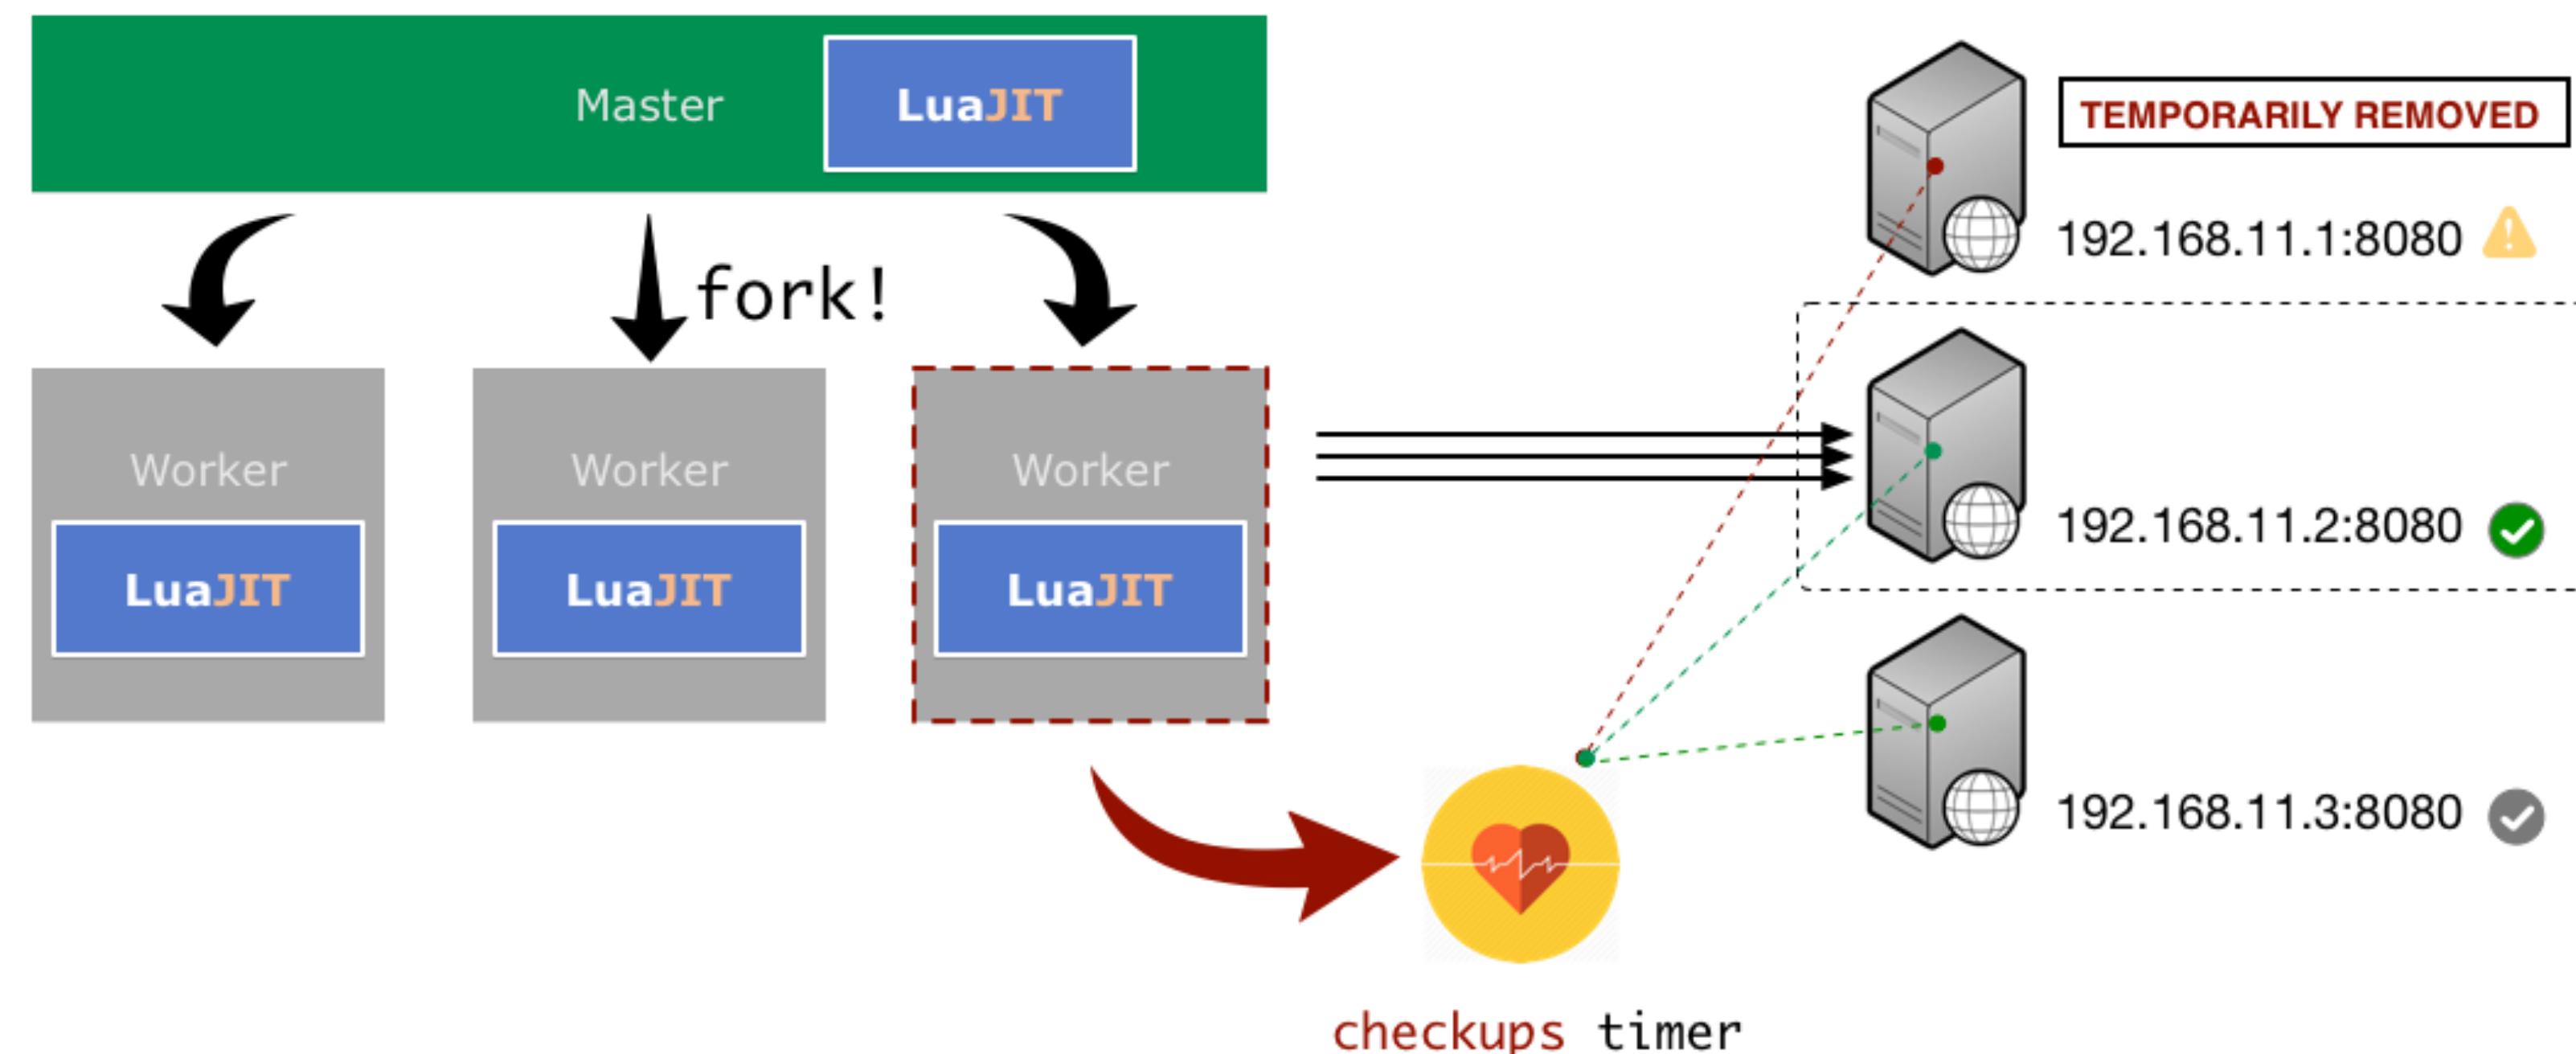
_M.global = {
    checkup_timer_interval = 5,
    checkup_timer_overtime = 60,
}

_M.api = {
    timeout = 2,
    typ = "general", -- http, redis, mysql etc.

    cluster = {
        { -- level 1
            try = 2,
            servers = {
                { host = "192.168.11.1", port = 8080, weight = 1 },
                { host = "192.168.11.2", port = 8080, weight = 2 },
            }
        },
        { -- level 2
            servers = {
                { host = "192.168.11.3", port = 8080, weight = 1 },
            }
        },
    },
}
```

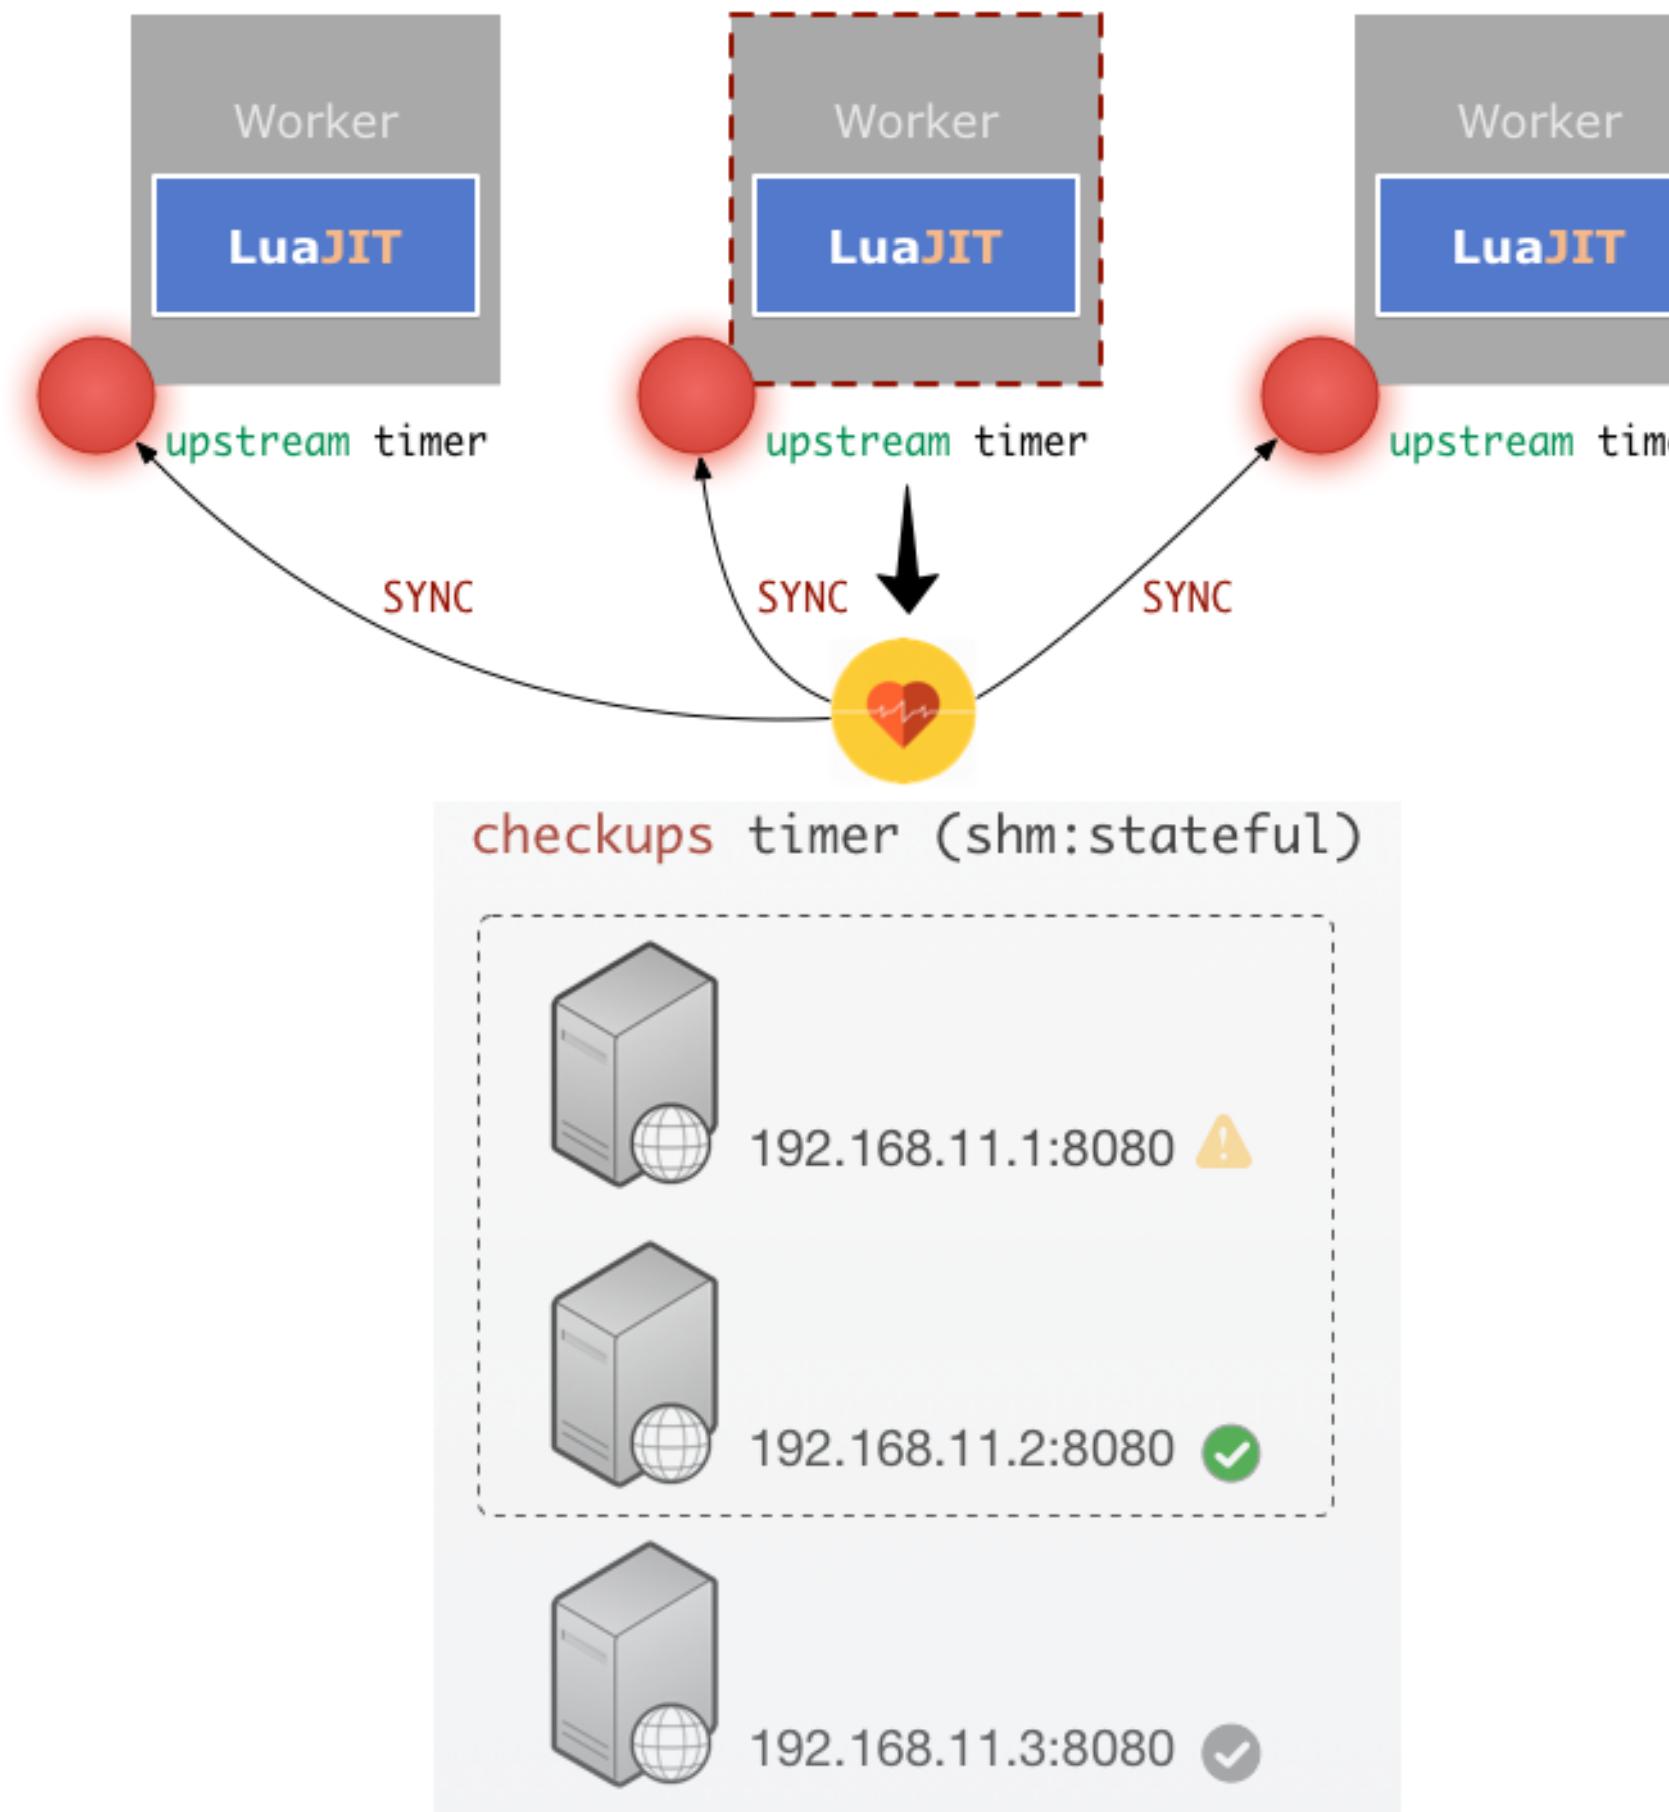


# Lua Upstream Health Checks: lua-resty-checkups



```
access_by_lua '  
    local checkups = require "resty.checkups"  
  
    -- only one timer is active among all the nginx workers  
    checkups.create_checker()  
';
```

# Lua Upstream Health Checks: checkups with nginx.conf



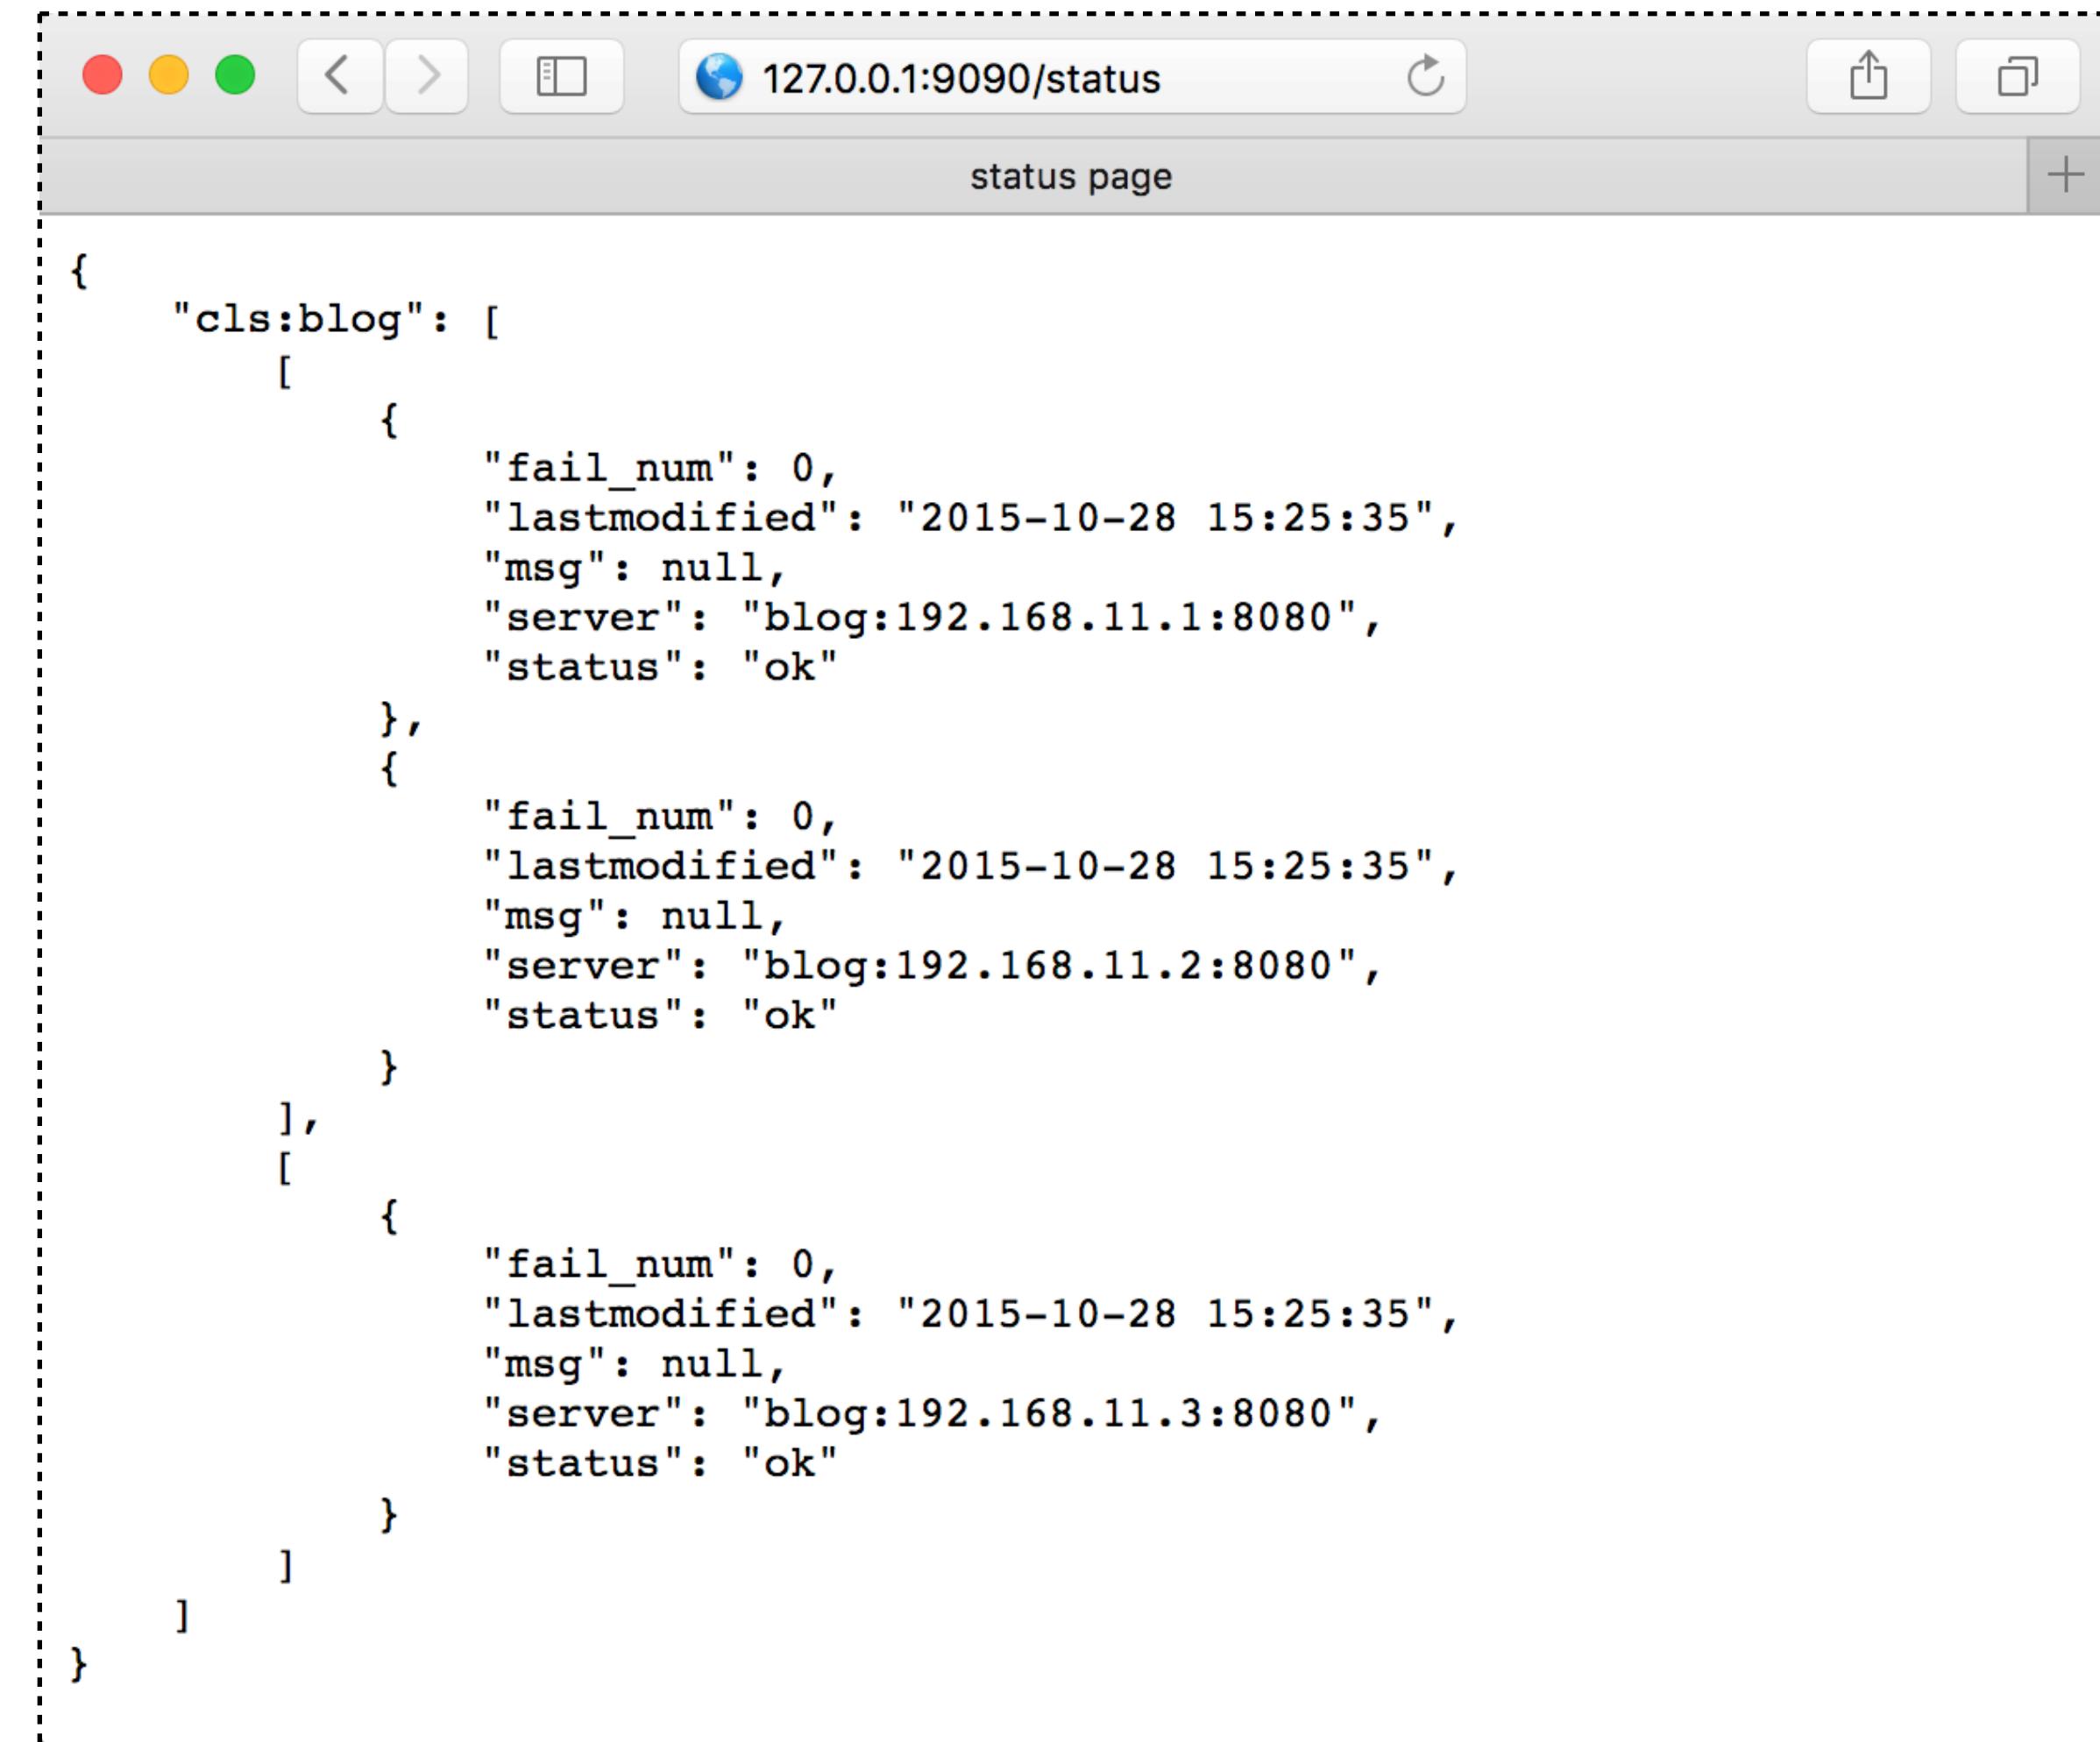
```
-- app/etc/config.lua

_M.global = {
    checkup_timer_interval = 5,
    checkup_timer_overtime = 60,
    ups_status_sync_enable = true,
    ups_status_timer_interval = 2,
}

_M.blog = {
    cluster = {
        { -- level 1
            try = 2,
            upstream = "blog.upyun.com",
        },
        { -- level 2
            upstream = "blog.upyun.com",
            upstream_only_backup = true,
        },
    },
}
```

lua-upstream-nginx-module

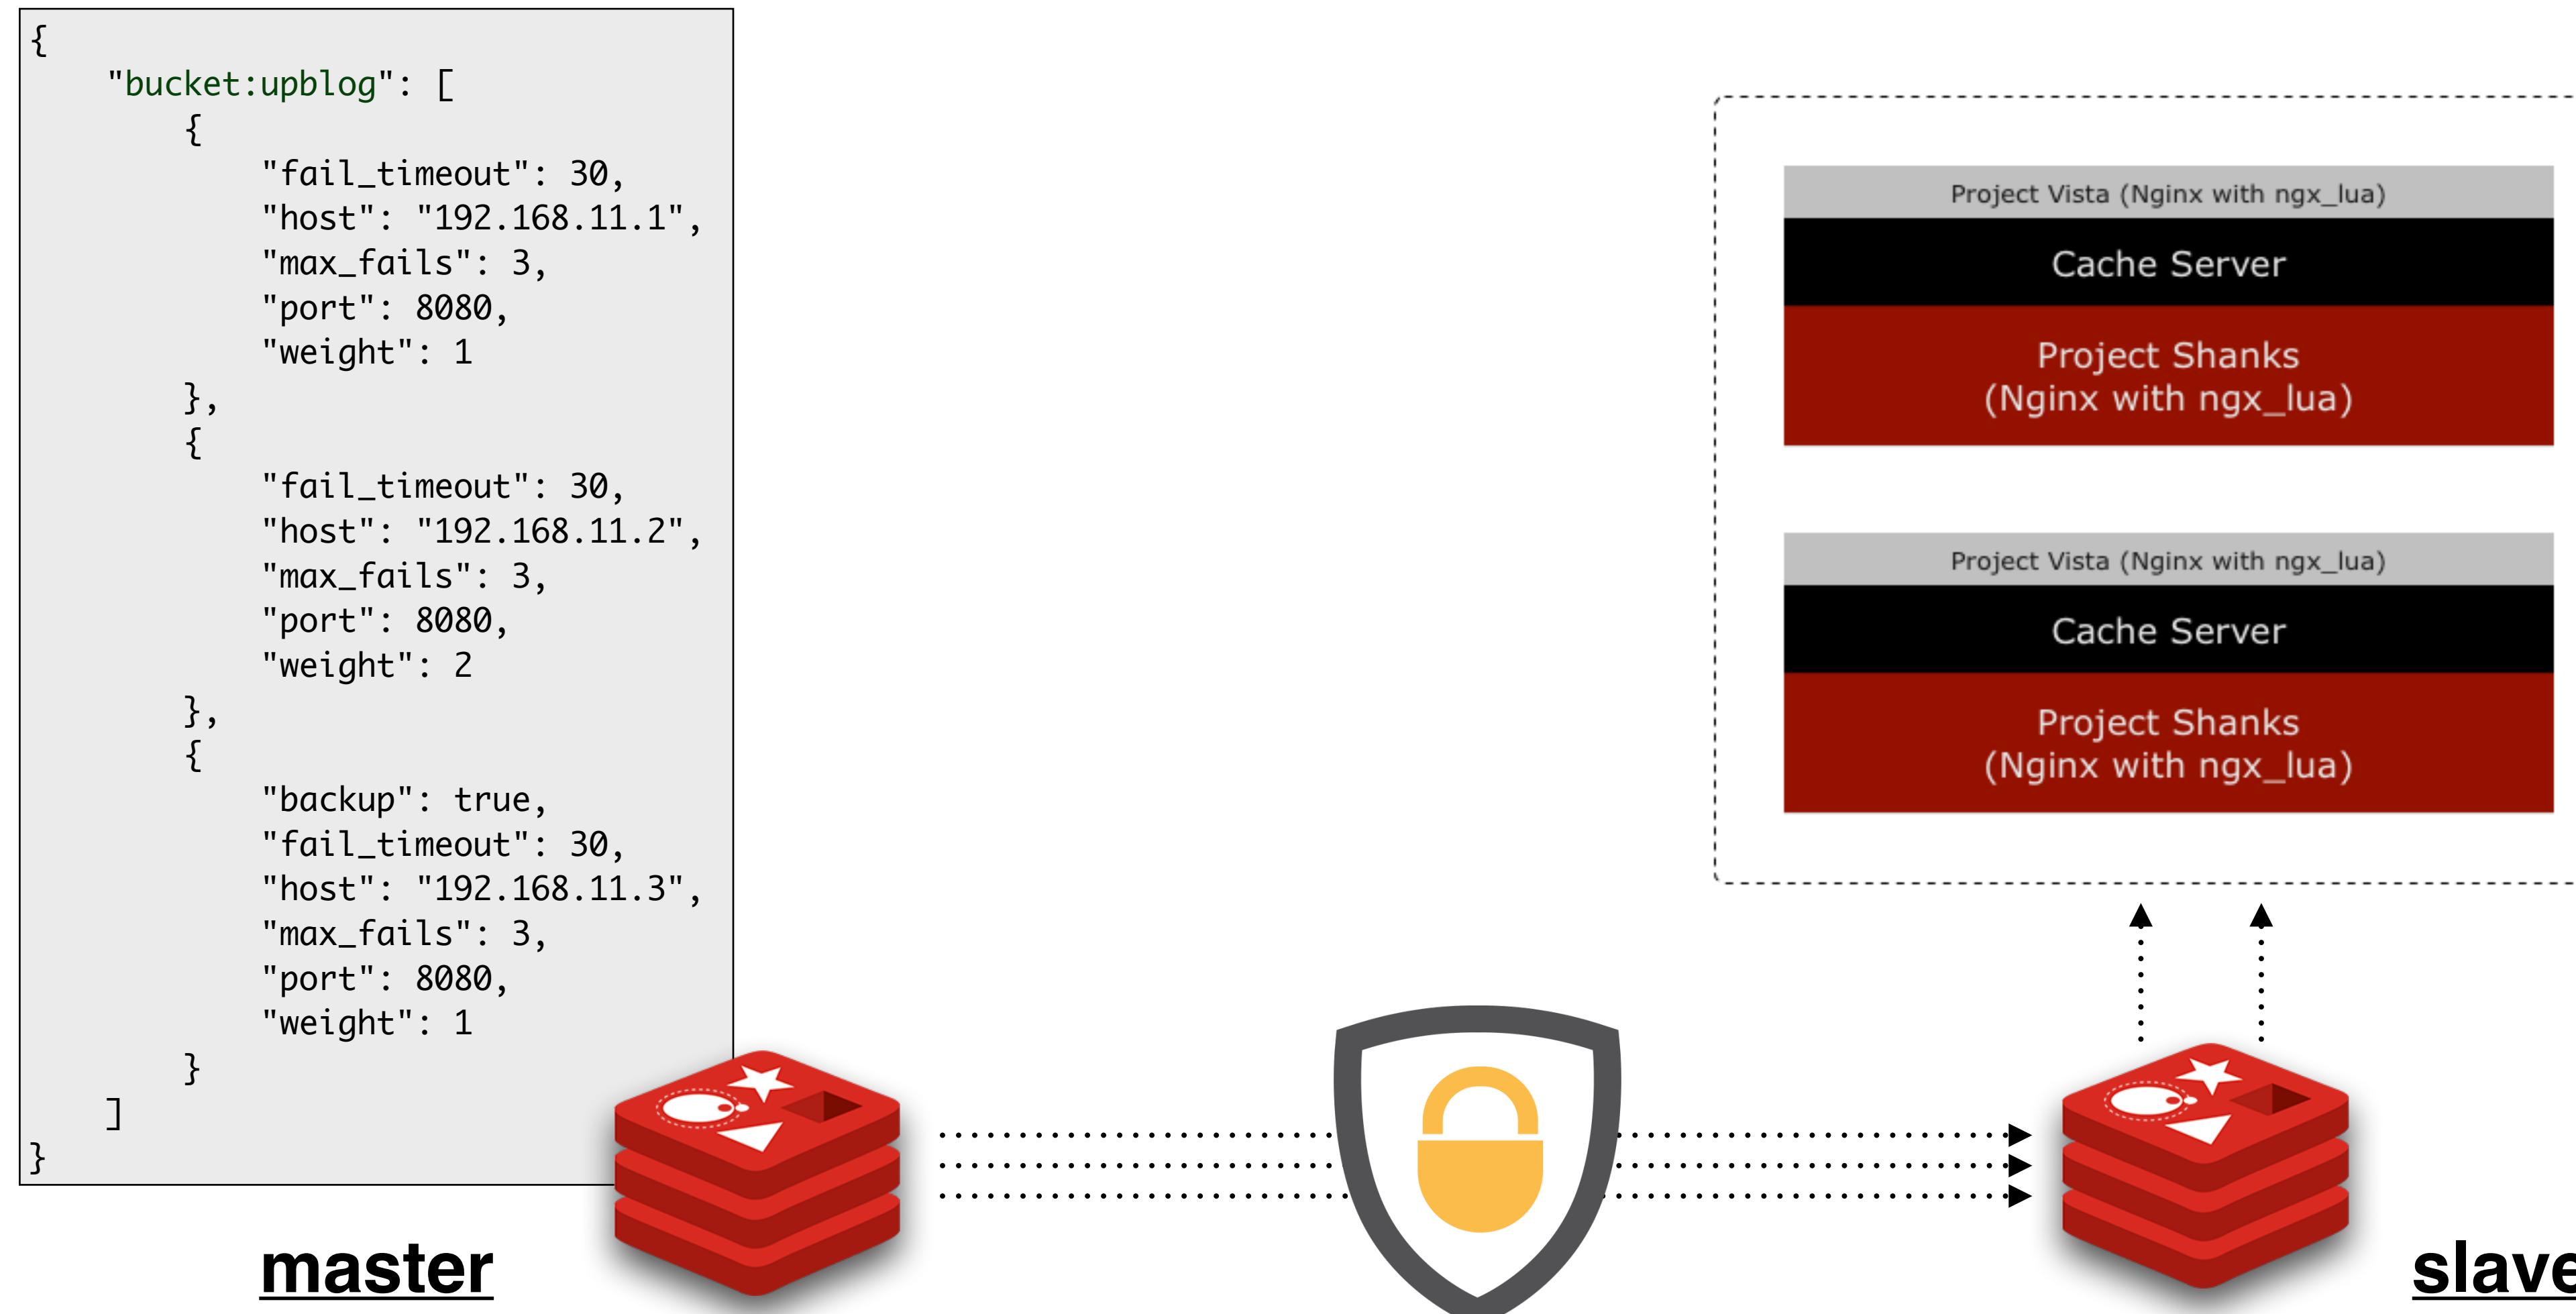
# Lua Upstream Health Checks: **checkups with status page**



A screenshot of a web browser window displaying a JSON status page. The window title is "status page" and the URL is "127.0.0.1:9090/status". The JSON output shows a configuration for the "cls:blog" class, which contains three upstream servers. Each server entry includes fields for fail\_num, lastmodified, msg, server, and status.

```
{
  "cls:blog": [
    [
      {
        "fail_num": 0,
        "lastmodified": "2015-10-28 15:25:35",
        "msg": null,
        "server": "blog:192.168.11.1:8080",
        "status": "ok"
      },
      {
        "fail_num": 0,
        "lastmodified": "2015-10-28 15:25:35",
        "msg": null,
        "server": "blog:192.168.11.2:8080",
        "status": "ok"
      }
    ],
    [
      {
        "fail_num": 0,
        "lastmodified": "2015-10-28 15:25:35",
        "msg": null,
        "server": "blog:192.168.11.3:8080",
        "status": "ok"
      }
    ]
  ]
}
```

# Lua Upstream Dynamically: Configure **Everything** as **JSON**



# Lua Metadata Cache:

## lua-resty-shcache

```
-- app/src/modules/metadata.lua

local shcache = require "resty.shcache"

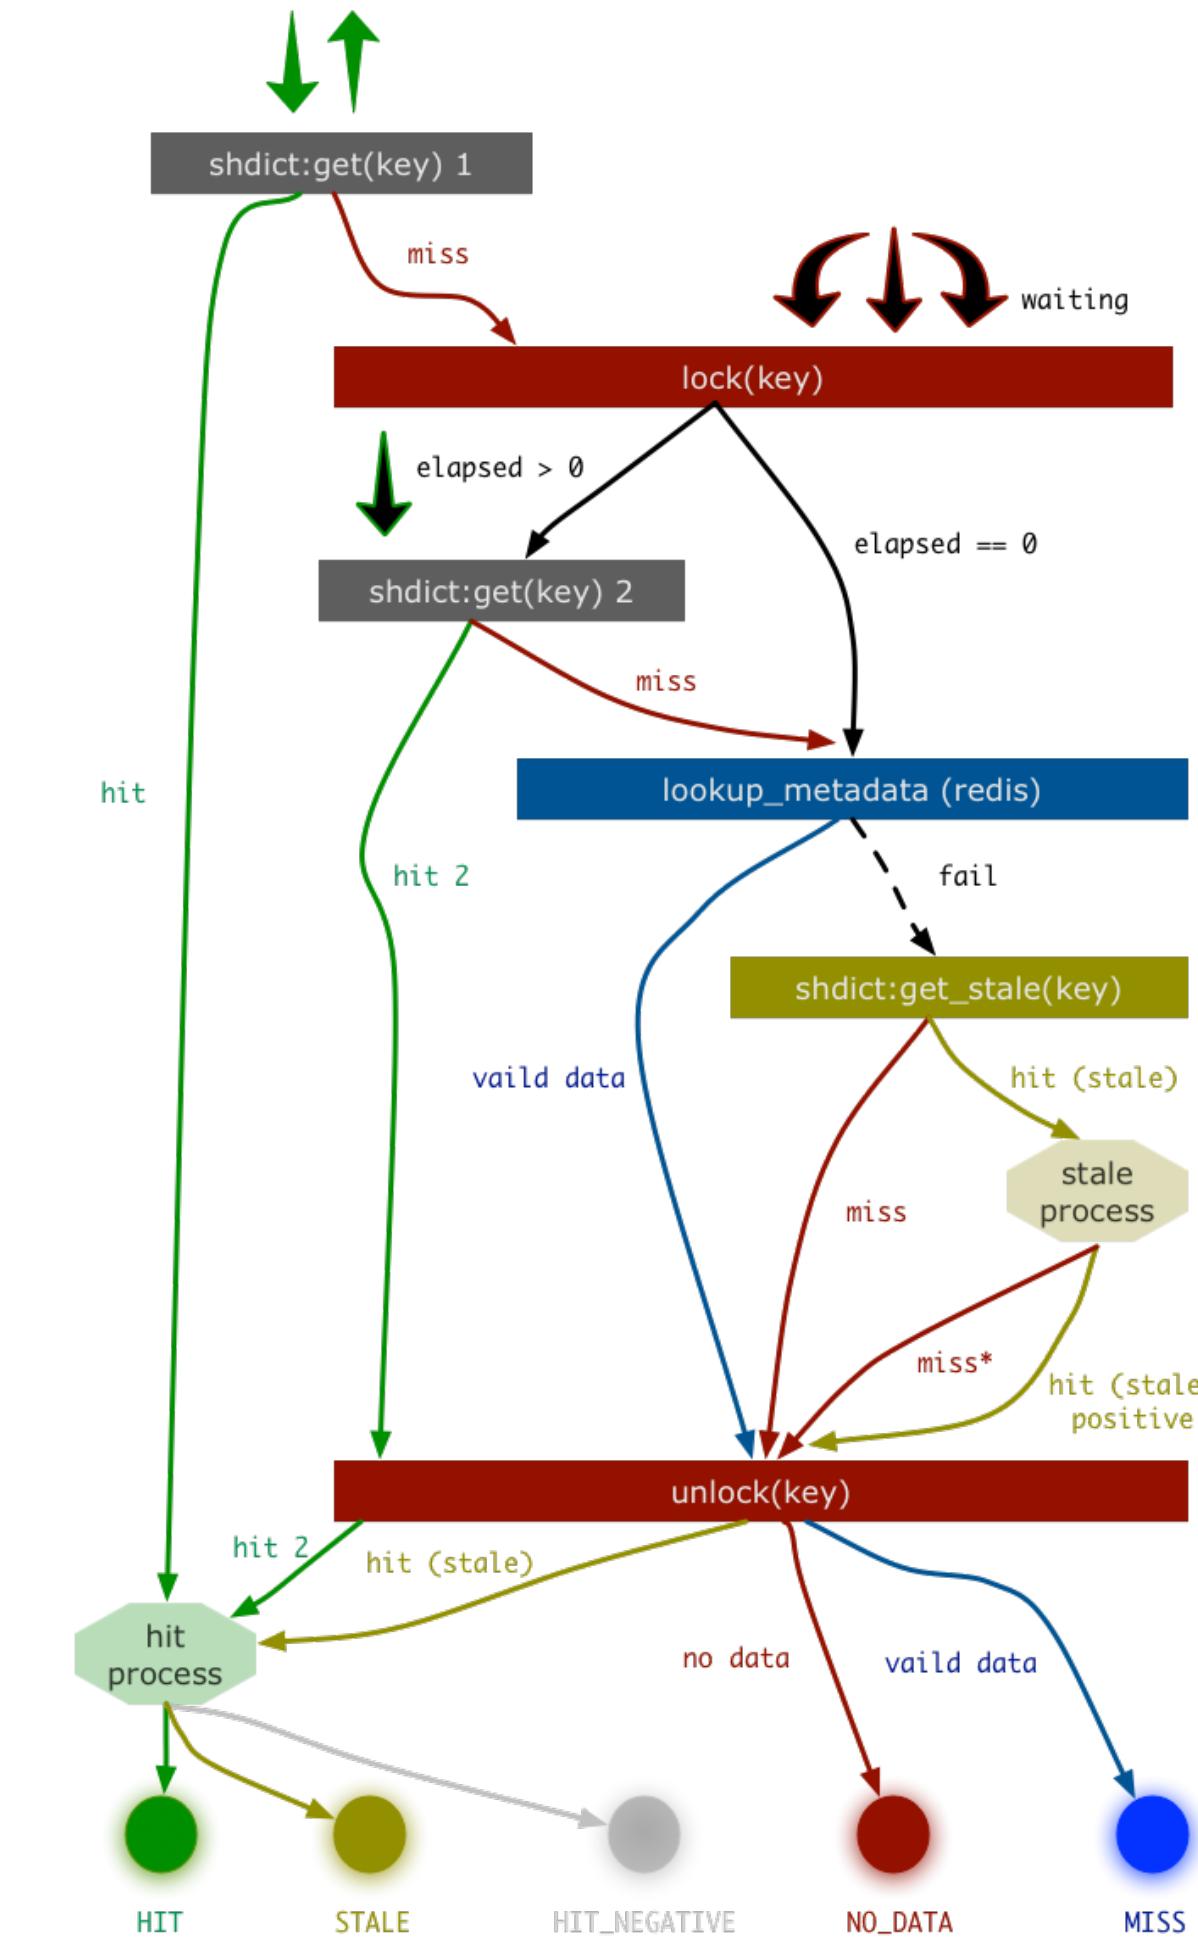
function _M.get_metadata(bucket)
    local lookup_metadata = function ()
        -- fetch from redis
        return res
    end

    local cache_data = shcache:new(
        ngx.shared.metadata,
        { external_lookup = lookup_metadata,
          encode = cmsgpack.pack,
          decode = cmsgpack.unpack,
        },
        { positive_ttl = cache_positive_ttl,
          negative_ttl = cache_negative_ttl,
          name = "metadata",
        }
    )

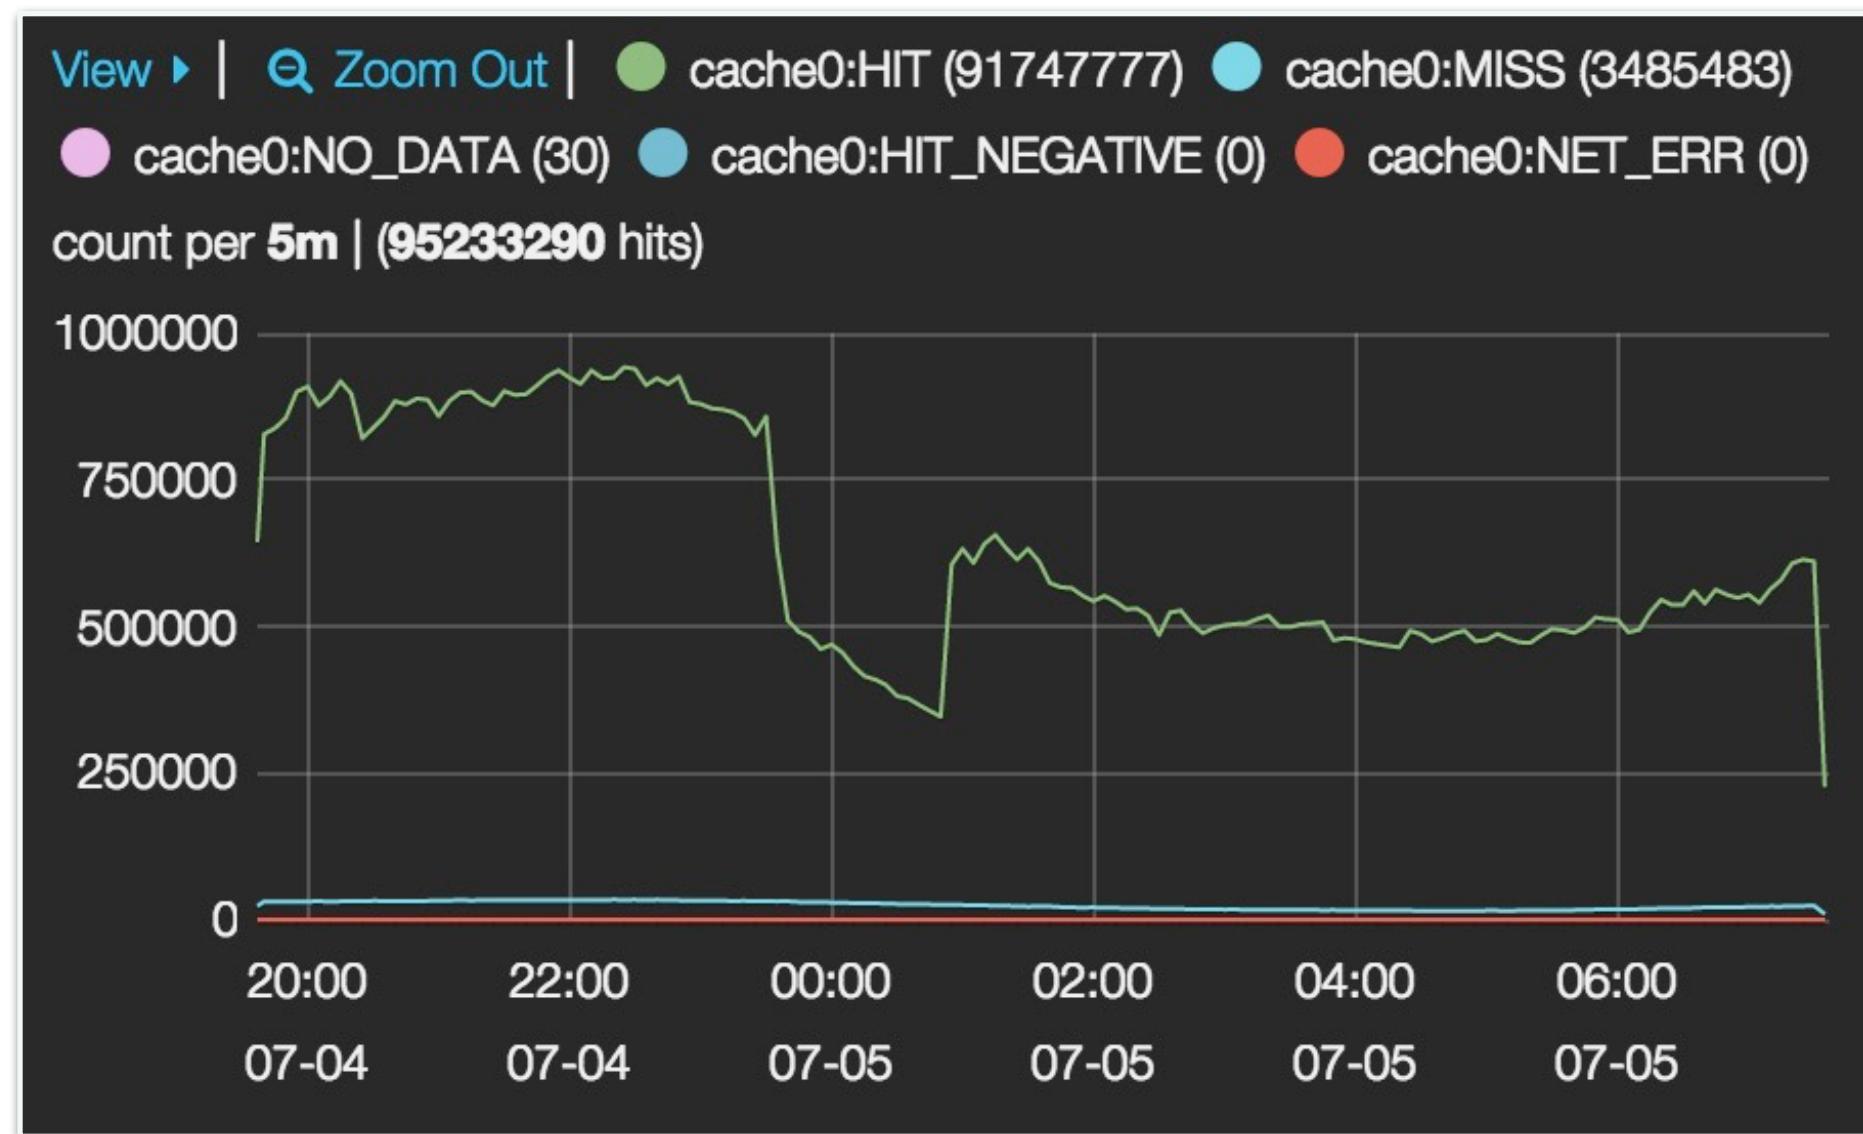
    -- local key = ...

    local data, _ = cache_data:load(key)
    if not data then
        return
    end

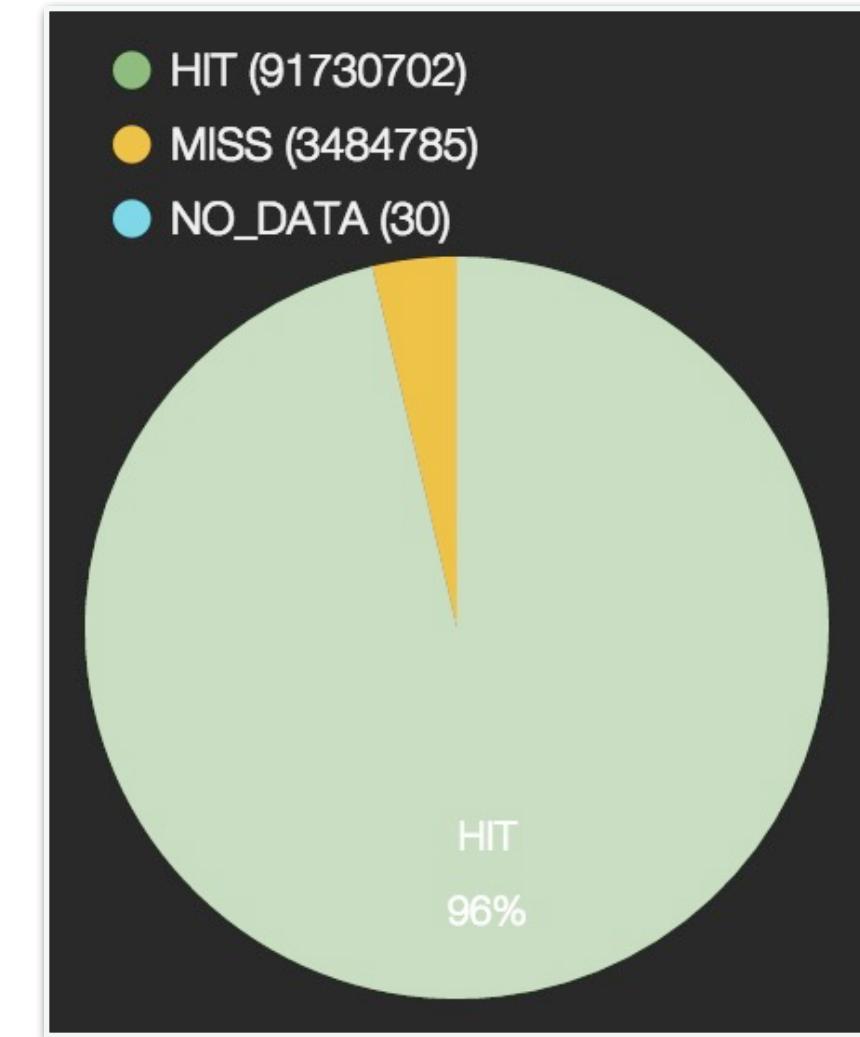
    return data
end
```



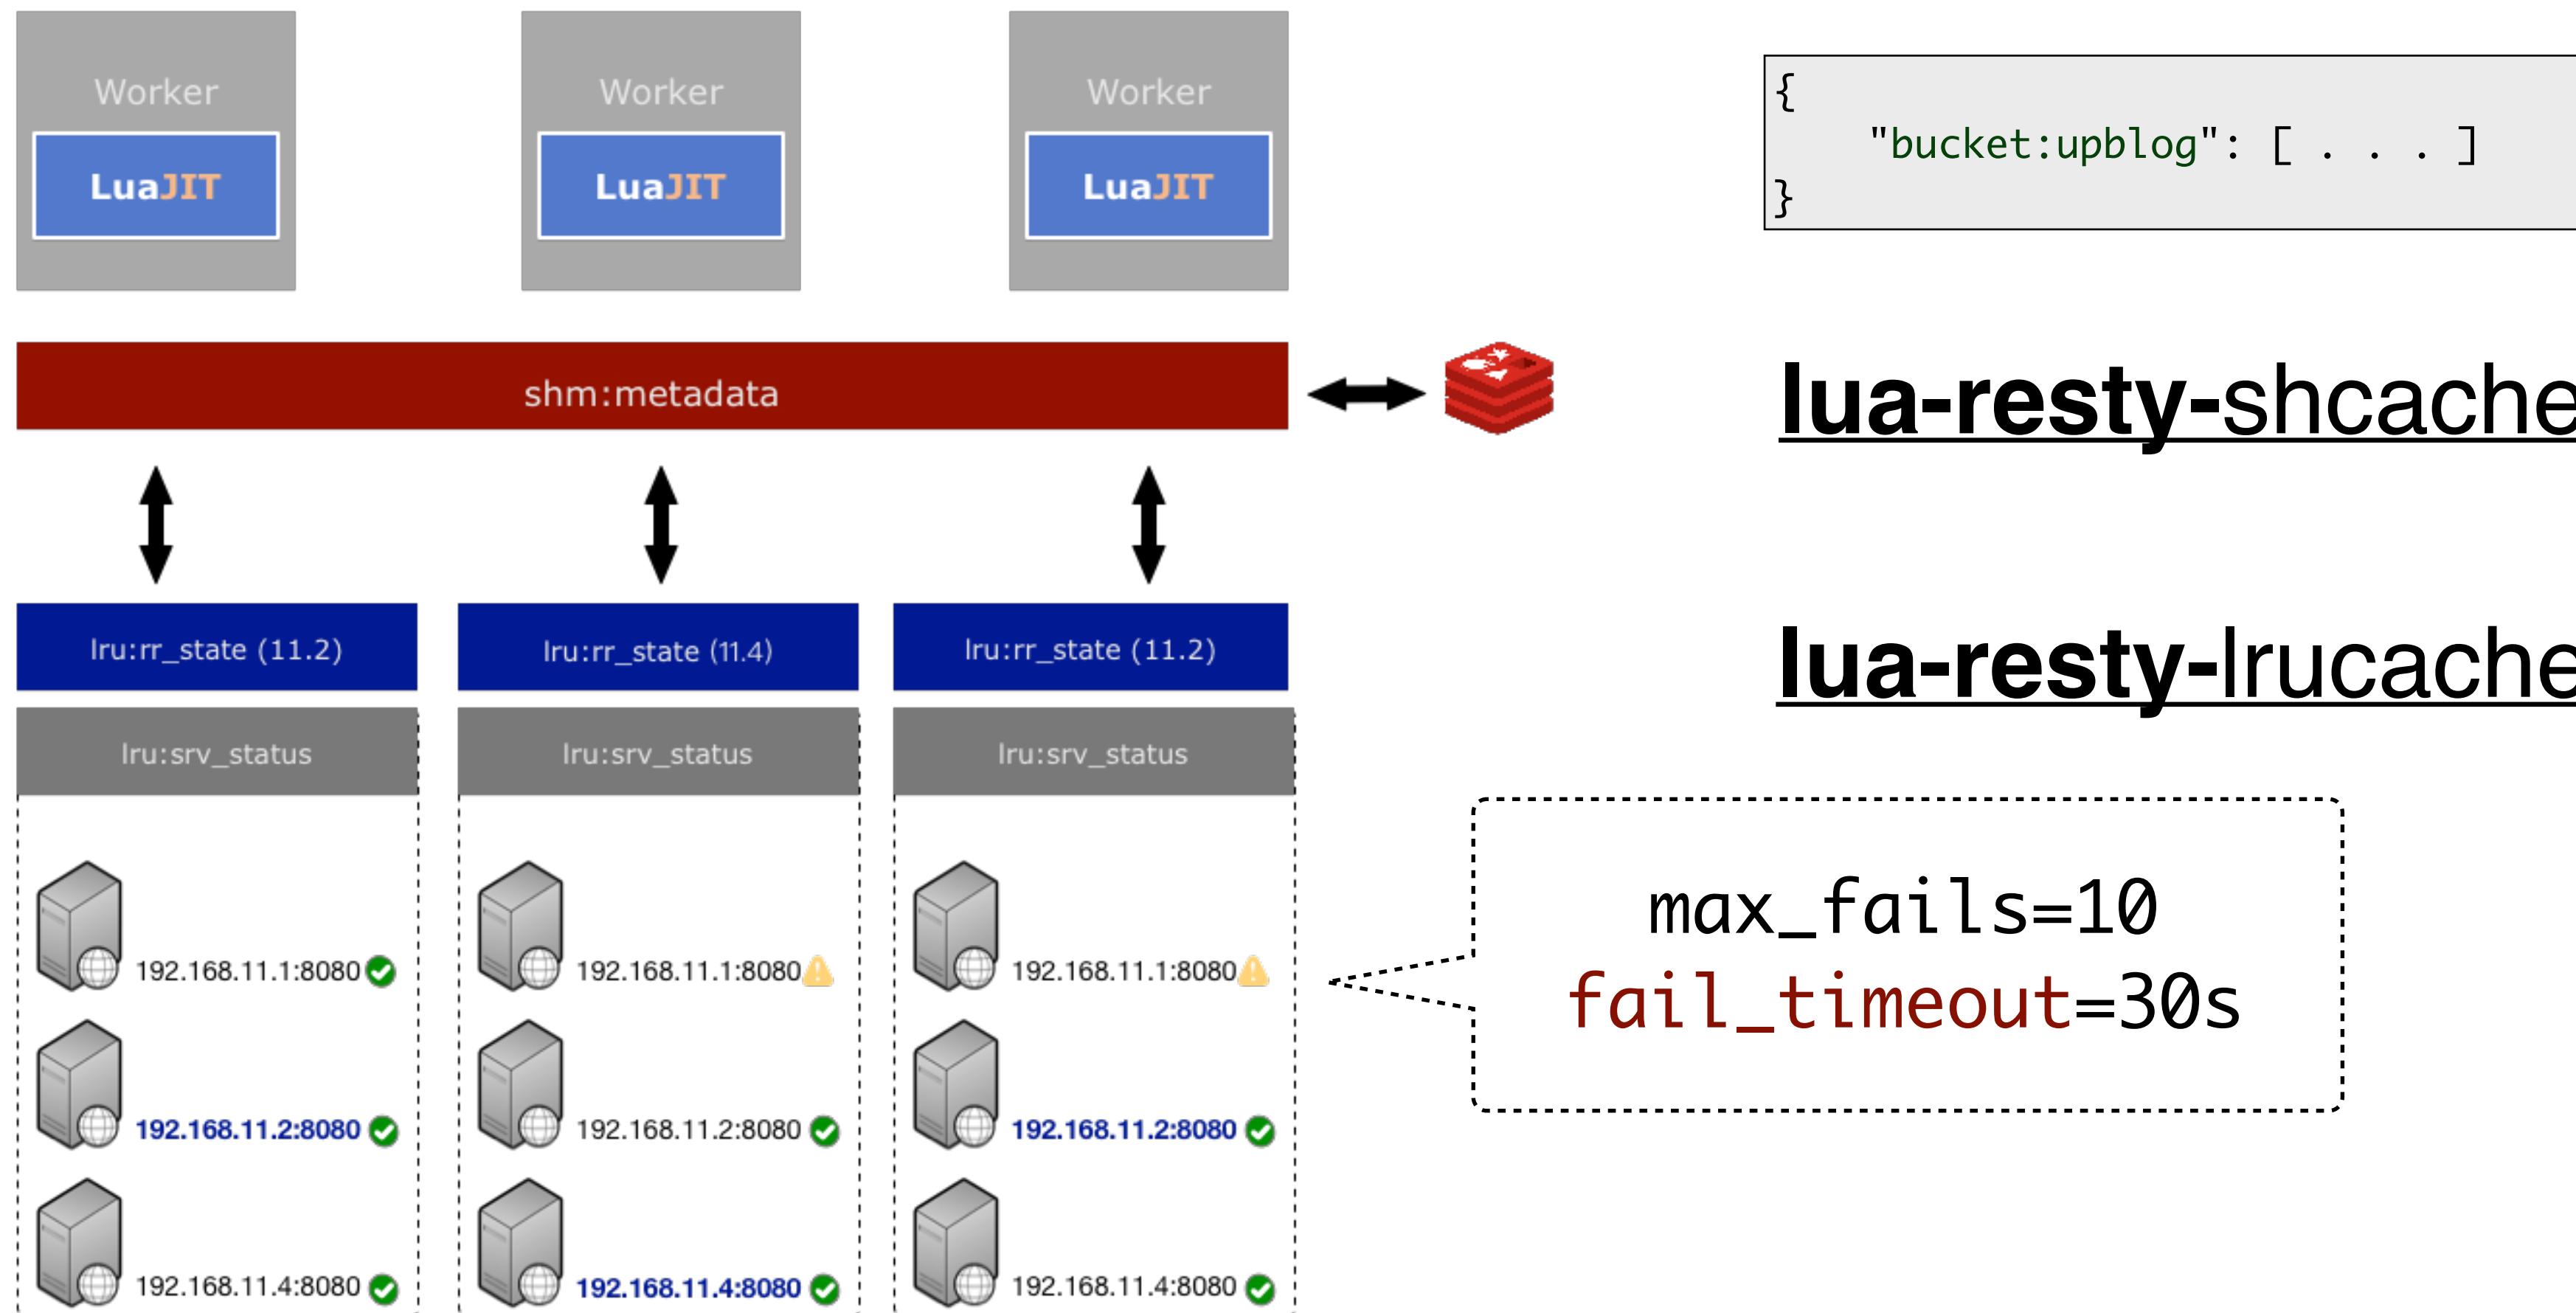
# Lua Metadata Cache: lua-resty-dbcache



- ★ **HIT**
- ★ **STALE**
- ★ **HIT\_NEGATIVE**
- ★ **NO\_DATA**
- ★ **MISS**
- ★ **??? (NET\_ERR)**



# Lua Upstream Dynamically: maintaining internal state



# Lua Upstream Load Balancing: round-robin with weight

```
function _M.reset_round_robin_state(cls)
    local rr = { index = 0, current_weight = 0 }
    rr.gcd, rr.max_weight, rr.weight_sum = _M.calc_gcd_weight(cls.servers)
    cls.rr = rr
end
```

```
cluster = {
    {
        servers = {
            { host = "127.0.0.1", port = 12351, weight = 1 },
            { host = "127.0.0.1", port = 12352, weight = 4 },
            { host = "127.0.0.1", port = 12353, weight = 3 },
            { host = "127.0.0.1", port = 12355, weight = 6 },
        }
    }
}
```

rr.index = 0  
rr.current\_weight = 0  
rr.gcd = 1  
rr.max\_weight = 6  
rr.weight\_sum = 14

# Lua Upstream Load Balancing: round-robin with weight

```
local try_servers_by_round_robin = function(cls, verify_server_status, callback)
```

```
local bad_servers = {}

for i = 1, #cls.servers, 1 do
    local srv, index, err = _M.select_round_robin_server(cls, verify_server_status, bad_servers)
    if not srv then
        return nil, err
    else
        local res, _ = callback(srv)

        if res then
            if srv.effective_weight ~= srv.weight then
                srv.effective_weight = srv.weight
                _M.reset_round_robin_state(cls)
            end
            return res
        end

        if srv.effective_weight > 1 then
            srv.effective_weight = floor(sqrt(srv.effective_weight))
            _M.reset_round_robin_state(cls)
        end

        bad_servers[index] = true
    end
end
```

# Lua Upstream Load Balancing: round-robin with weight

```
function _M.select_round_robin_server(cls, verify_server_status, bad_servers)
```

```
local rr = cls.rr
```

```
local servers = cls.servers
```

```
local index = rr.index
```

```
local current_weight = rr.current_weight
```

```
local gcd = rr.gcd
```

```
local max_weight = rr.max_weight
```

```
local weight_sum = rr.weight_sum
```

```
local failed = 1
```

```
repeat
```

```
until failed > weight_sum
```

**TALK IS  
CHEAP**

# Lua Upstream Load Balancing: round-robin with weight

```
index = index % #servers + 1
if index == 1 then
    current_weight = current_weight - gcd
    if current_weight <= 0 then current_weight = max_weight end
end

local srv = servers[index]
if srv.effective_weight >= current_weight then
    cls.rr.index, cls.rr.current_weight = index, current_weight
    if not bad_servers[index] then
        if verify_server_status then
            if verify_server_status(srv) then
                return srv, index
            else
                if srv.effective_weight > 1 then
                    srv.effective_weight, index, current_weight, failed_count = 1, 0, 0, 0
                    _M.reset_round_robin_state(cls)
                    gcd, max_weight, weight_sum = cls.rr.gcd, cls.rr.max_weight, cls.rr.weight_sum
                end
                failed = failed + 1
            end
        else
            return srv, index
        end
    else
        failed = failed + 1
    end
end
end
```

**repeat** ..... **until** failed > weight\_sum

# Lua Upstream Load Balancing: round-robin with weight

```
local verify_server_status = function(srv)
    local peer_key = _gen_key(srv)

    local peer_status = cjson.decode(state:get(PEER_STATUS_PREFIX .. peer_key))
    if peer_status == nil or peer_status.status ~= _M.STATUS_ERR then
        return true
    end

    return
end
```

★**STATUS\_OK** = 0

★**STATUS\_UNSTABLE** = 1

★**STATUS\_ERR** = 2

# Lua Upstream Load Balancing:

## ==== TEST 1: round-robin single level

```
-- http_config eval
"$::HttpConfig" . "$::InitConfig"
-- config
location = /t {
    content_by_lua '
        local checkups = require "resty.checkups"
        checkups.create_checker()
        ngx.sleep(2)
        local dict = {
            [12351] = "A",
            [12352] = "B",
            [12353] = "C",
            [12355] = "E",
        }
        local cb_ok = function(srv)
            ngx.print(dict[srv.port])
            return 1
        end

        for i = 1, 30, 1 do
            local ok, err = checkups.ready_ok("single_level", cb_ok)
            if err then
                ngx.say(err)
            end
        end
    ';
}
-- request
GET /t
-- response_body: EEBEBCEBCEABCEEEBEBCEBCEABCEEE
```

```
_M.single_level = {
    cluster = {
        {
            servers = {
                { host = "127.0.0.1", port = 12351, weight = 1 },
                { host = "127.0.0.1", port = 12352, weight = 4 },
                { host = "127.0.0.1", port = 12353, weight = 3 },
                { host = "127.0.0.1", port = 12355, weight = 6 },
            }
        }
    }
}
```

**EEBEBCEBCEABCE**

.....

# Lua Upstream Load Balancing: consistent-hash and more

```
cluster = {
  {
    servers = {
      { host = "127.0.0.1", port = 12351, weight = 1 },
      { host = "127.0.0.1", port = 12352, weight = 4 },
      { host = "127.0.0.1", port = 12353, weight = 3 },
      { host = "127.0.0.1", port = 12355, weight = 6 },
    }
  },
  {
    servers = {
      { host = "127.0.0.1", port = 12354, weight = 1 },
      { host = "127.0.0.1", port = 12356, weight = 2 },
    }
  }
}
```

★**try\_servers\_by\_round\_robin**  
★**try\_cluster\_by\_round\_robin**

★**try\_servers\_by\_consistent\_hash**  
★**try\_cluster\_by\_consistent\_hash**

## CDN 设置

X

\* 回源 Host:  域名跟随  自定义

回源方式:  HTTP 协议回源  HTTPS 协议回源  协议跟随

\* 源站线路:  电信  移动  联通  BGP  其他

## 电信

+

回源地址	端口号	线路属性	轮询权重	最大失败次数	静默时间(秒)
192.168.11.1	: 80	主线路	1	10	30
192.168.11.2	: 80	主线路	2	10	30
192.168.11.3	: 80	备用线路	1	10	30

## 联通

+

回源地址	端口号	线路属性	轮询权重	最大失败次数	静默时间(秒)
192.168.12.1	: 80	主线路	1	10	30

取消

确定



tianchaijz:

**"\$WHEN(\$MATCH(\$\_URI, '^/foo/.\*'))\$ADD\_REQ\_HEADER(X-Foo, bar)"**



Marco: I GOT IT !

**Edge Server**

# Lua Custom URL rewrite: lua-resty-rewrite | variables

**`$_HOST`**

**`$_SCHEME`**

**`$_METHOD`**

**`$_SYM_sym`**

**`$_HOST_n`**

**`$_POST_name`**

**`$_URI`**

**`$_HEADER_name`**

**`$_GET_name`**

**`$_COOKIE_name`**

**`$_RANDOM`**

**`$_QUERY`**

**`$_RANDOM_n`**

# Lua Custom URL rewrite: lua-resty-rewrite I functions

\$ENCODE_BASE64(E)	<b>\$UPPER(E)</b>
	\$ALL(E1, E2, ...)
\$ANY(E1, E2, ...)	\$DECODE_BASE64(E)
	<b>\$LOWER(E)</b>
	\$WHEN(E1, E2, ...)
<b>\$SUB(E1, from, to)</b>	\$PCALL(E)
\$GT(E1, E2)	\$ADD_REQ_HEADER(E1, E2)
\$GE(E1, E2)	\$DEL_REQ_HEADER(E1)
\$EQ(E1, E2)	\$ADD_RSP_HEADER(E1, E2)

# Lua Custom URL rewrite: lua-resty-rewrite | break

```
rewrite /download/(.*)/(.*) /$1/$2.mp3?_session=$_COOKIE_id?  
rewrite /download/(.*)/(.*) /$1/$2.mp3?user=$_HOST_1 break
```

...



See More: <https://github.com/upyun/docs/issues/5>

**http://io.upyun.com/2015/03/09/hello-world/?foo=bar**

**[scheme] [host] [path] [query]**

# Lua Custom Cache-Control:

## Using specific URI rules

```
location ^~ /www/ {  
    if ($query_string ~* "foo=bar") {  
        expires 300s;  
    }  
}  
  
location ^~ /images/ {  
    expires 1h;  
}  
  
location ~* \.jpg$ {  
    expires 1d;  
}
```

特殊缓存内容		不缓存内容	编辑	删除所选项
<input type="checkbox"/>	全选	缓存规则	缓存时间 (秒) <input type="text" value="?"/>	
<input type="checkbox"/>		/www/*?foo=bar	300	
<input type="checkbox"/>		/images/*	3600	
<input type="checkbox"/>		*.jpg	86400	

# Lua Custom SSL: Certificates Load & OCSP stapling

HTTPS 服务方式: 默认 UPYUN 域名      你可以开启自主域名的 HTTPS 服务, 立即购买      [添加 SSL 证书](#)

证书编号	证书颁发对象	使用组织名称	有效期	已配置域名	操作
02af75eee15a0266d59a48d0f34e1a9d	www.sw.com	浙江季产品网络集团	2015-06-10 - 2015-07-10	0 个	<a href="#">管理</a> <a href="#">删除</a> <a href="#">查看</a>
6128f9efa587cc20ab16ea69b1b0e5b6	www.sw.com	浙江季产品网络集团	2015-06-01 - 2015-07-01	0 个	<a href="#">管理</a> <a href="#">删除</a> <a href="#">查看</a>
UPYUN 默认 HTTPS 证书					9 个 <a href="#">管理</a>

使用说明

1. 一个绑定域名只能使用一个 SSL 证书, 配置开启 HTTPS 服务;
2. 泛域名证书配置给子域名开启 HTTPS 服务, 需到对应空间下的“通用-域名管理”操作;
3. 默认 UPYUN 域名的 HTTPS 服务使用, 按照空间进行配置管理;
4. HTTPS 服务功能配置生效时间, 全网 1~10 分钟。

```
server {
  listen 443 ssl;
  server_name upyun.com;

  ssl_certificate      upyun.com.pem;
  ssl_certificate_key  upyun.com.key;

  ssl_stapling on;
  ssl_stapling_verify on;
  ssl_trusted_certificate /etc/ssl/private/ca-certs.pem;
}
```

# Lua Custom Logging:

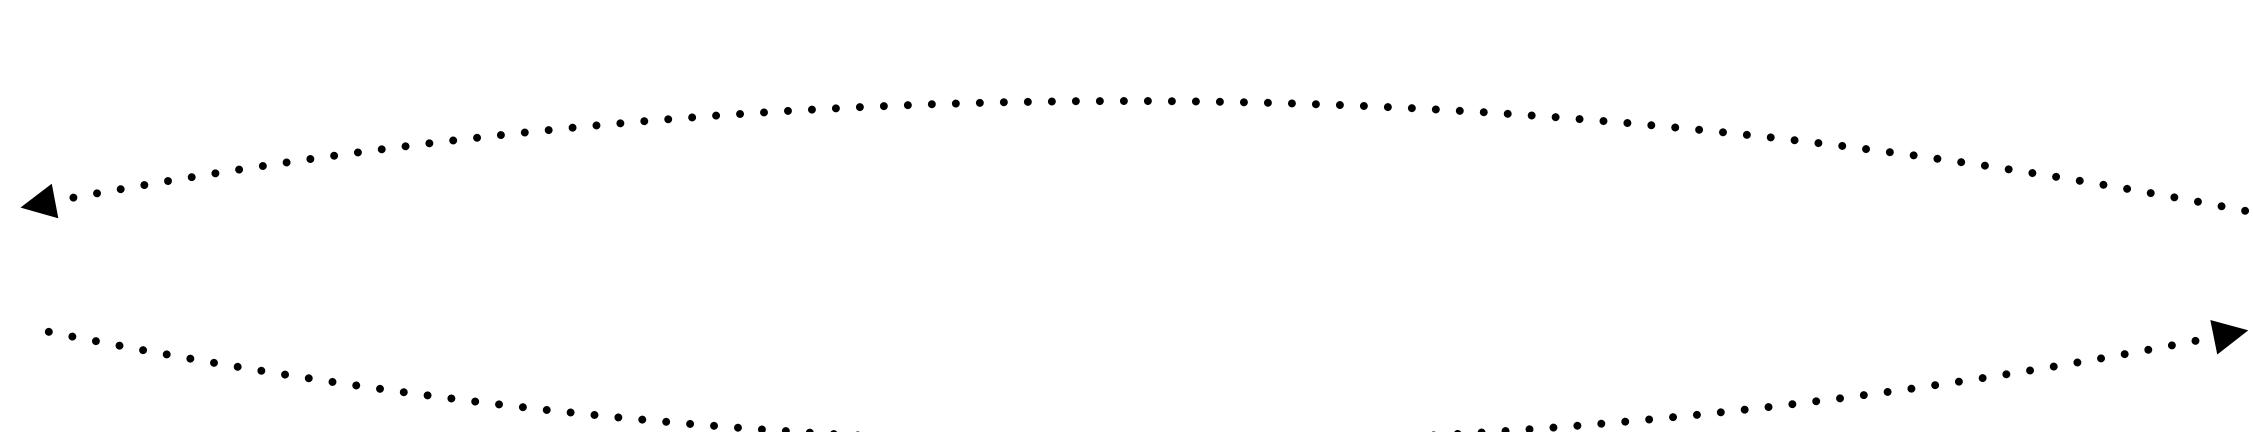
## lua-resty-logger-socket

```
log_format combined '$remote_addr - $remote_user [$time_local] '\n        '"$request" $status $body_bytes_sent '\n        '"$http_referer" "$http_user_agent"';\n\nserver {\n    access_log /path/to/access.log combined buffer=4096;\n    . . .\n}
```

Edge Server



**bucket:hbimg = {"enable":true,"ratio":0.1}**

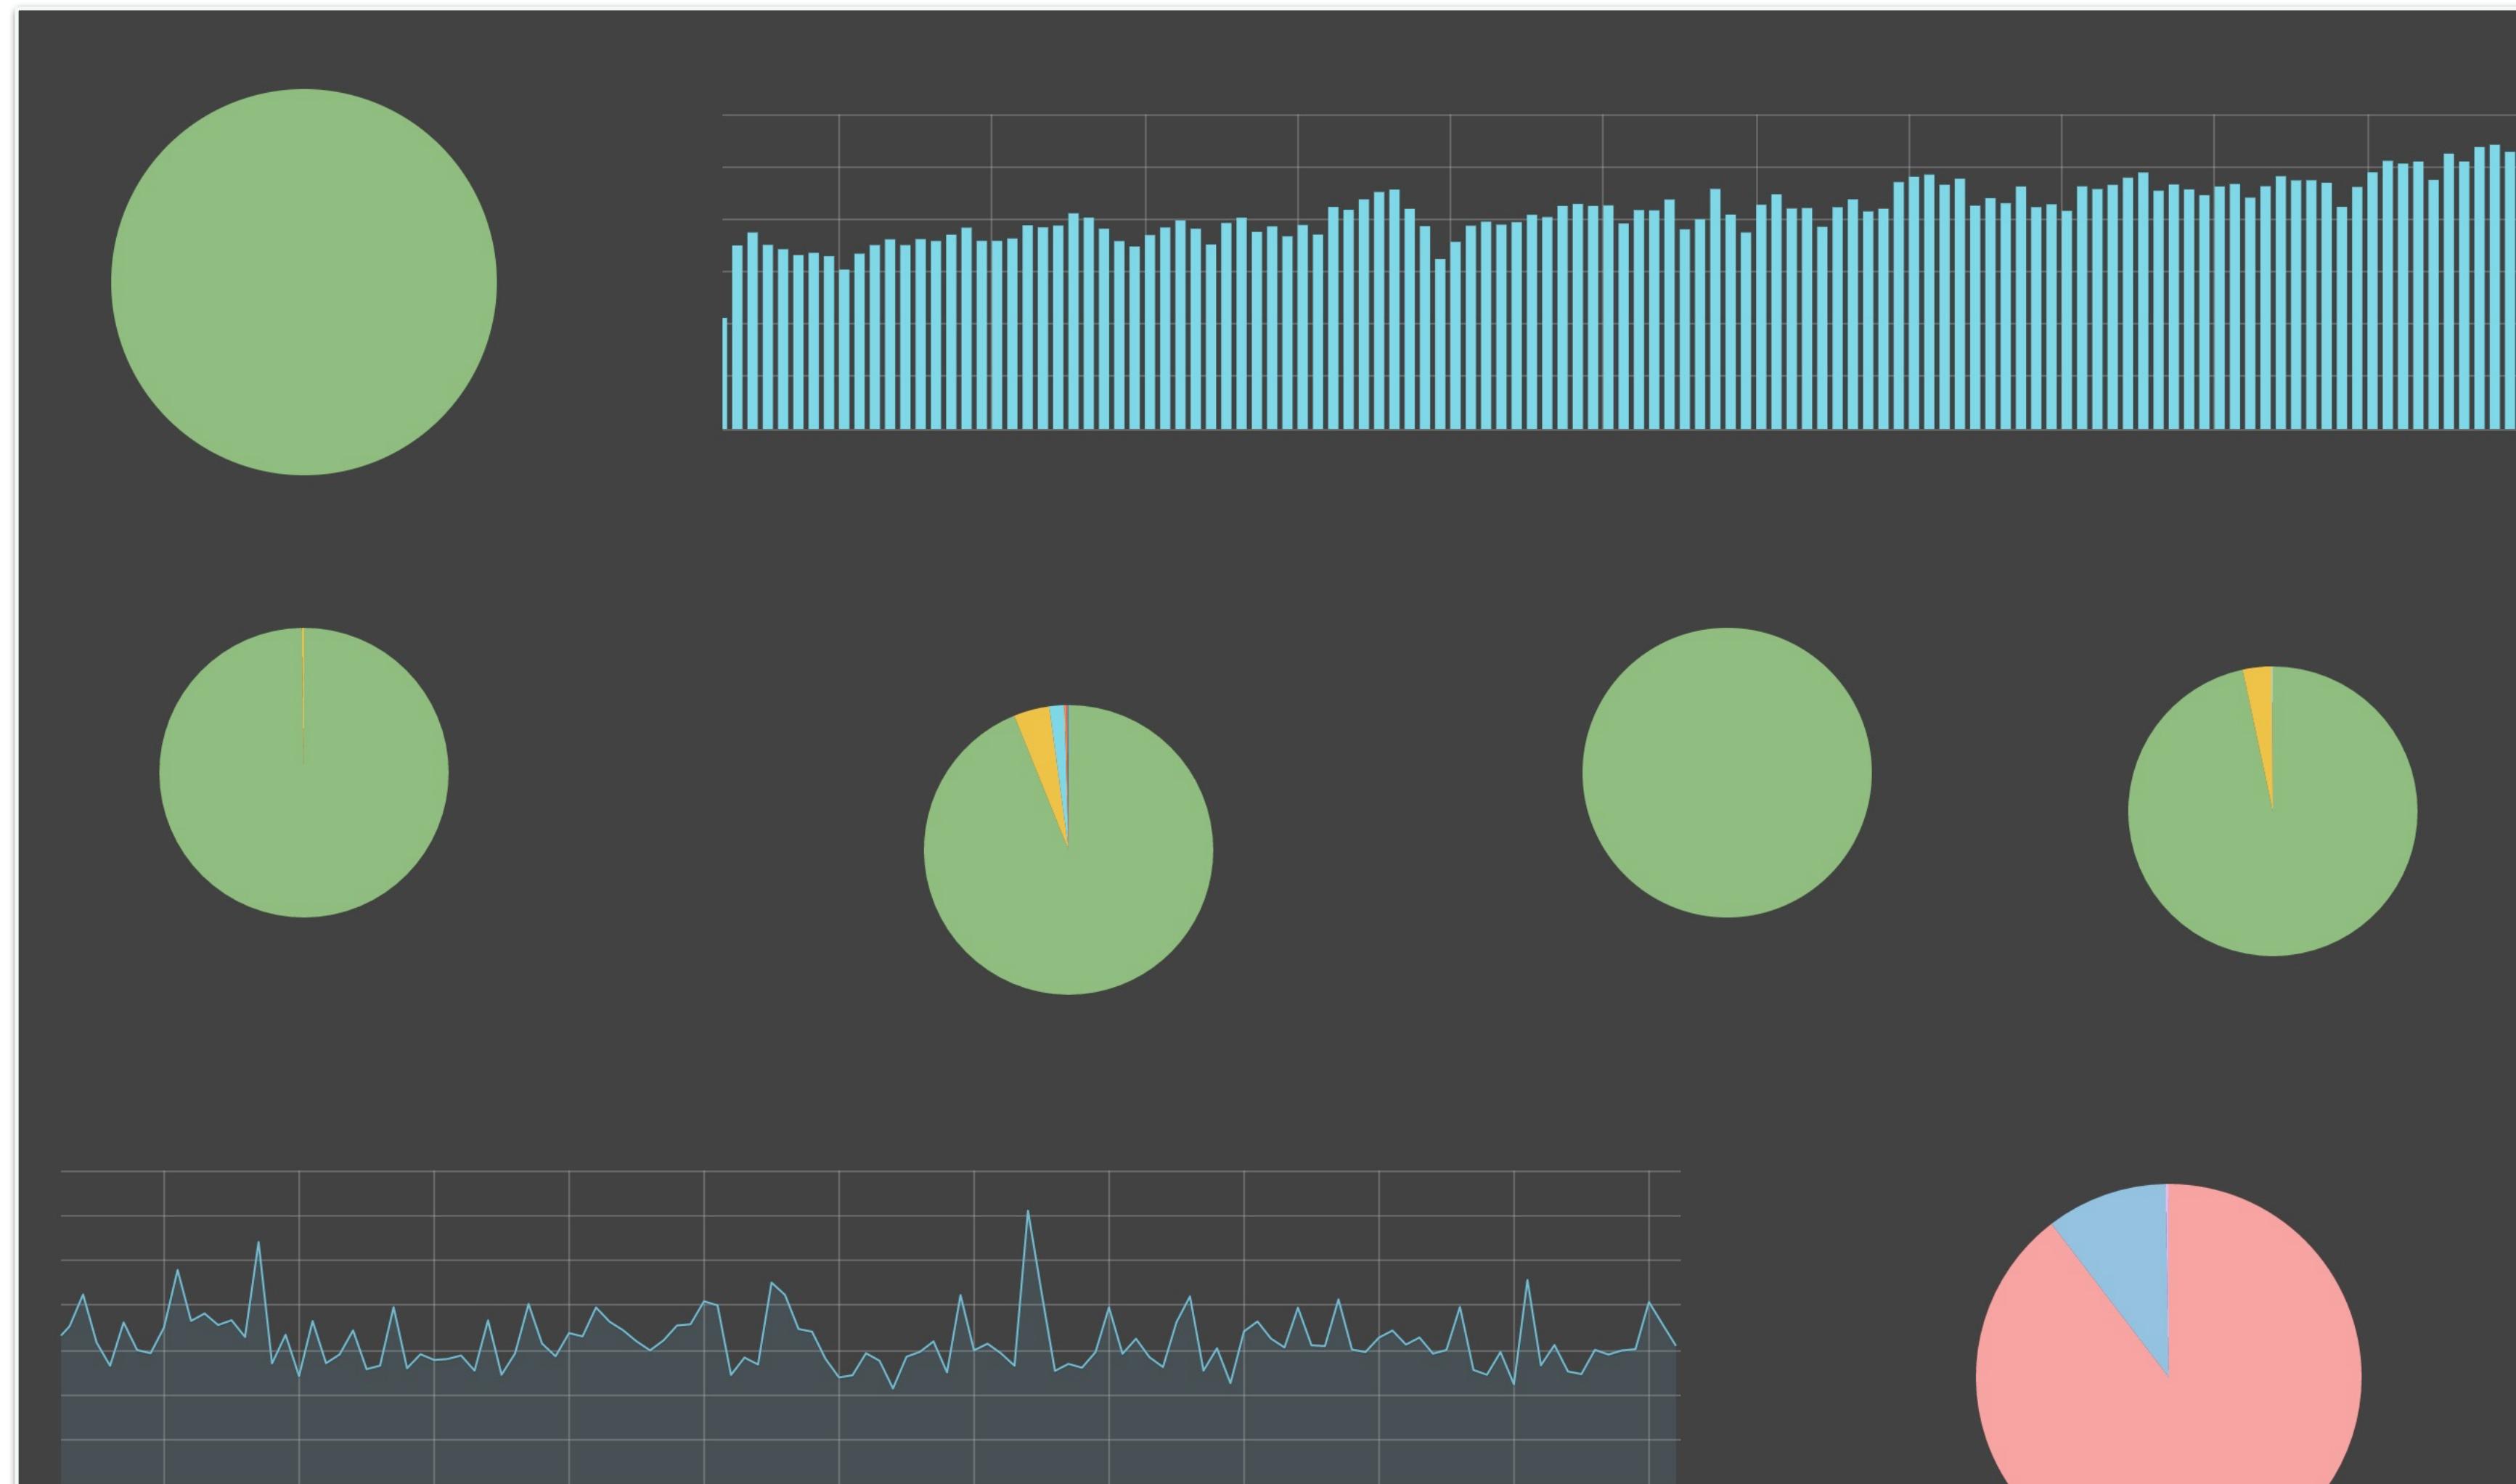


**logger.log(cjson.encode(log\_msg\_table) .. "\n")**

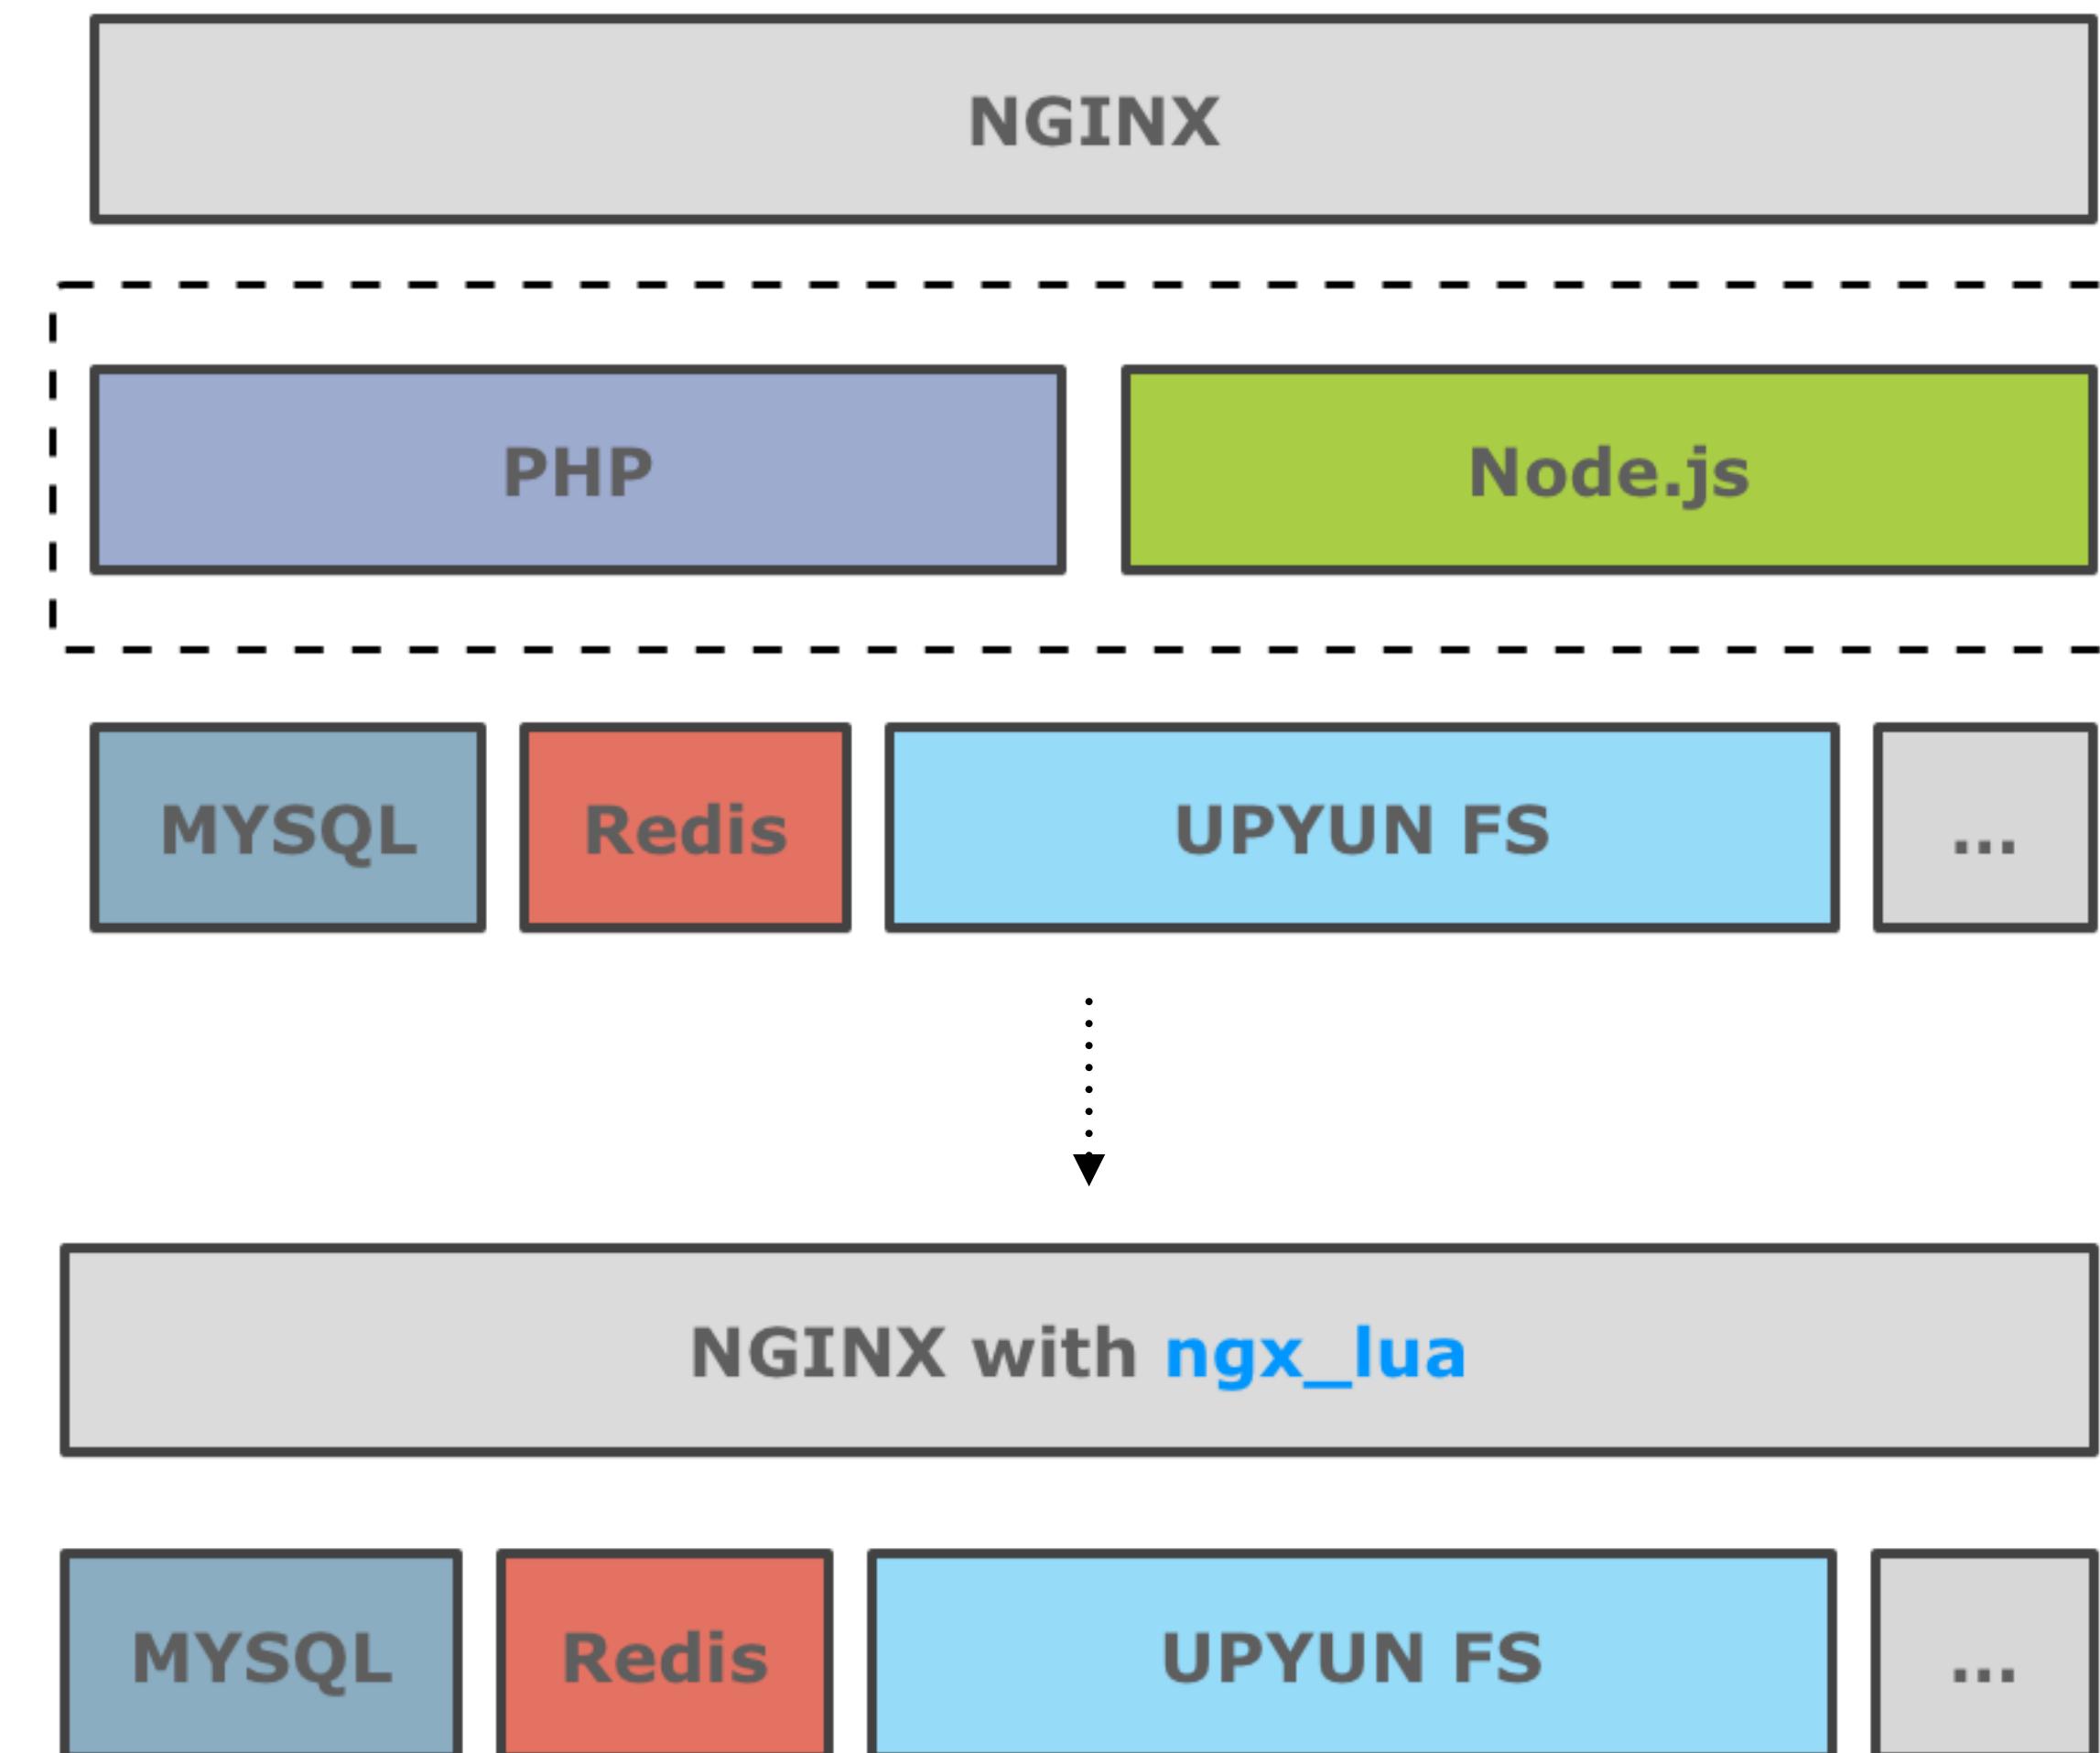
UPYUN LOG

# UPYUN LOG Platform:

## HAProxy + Heka + Kafka + Elasticsearch + Kibana

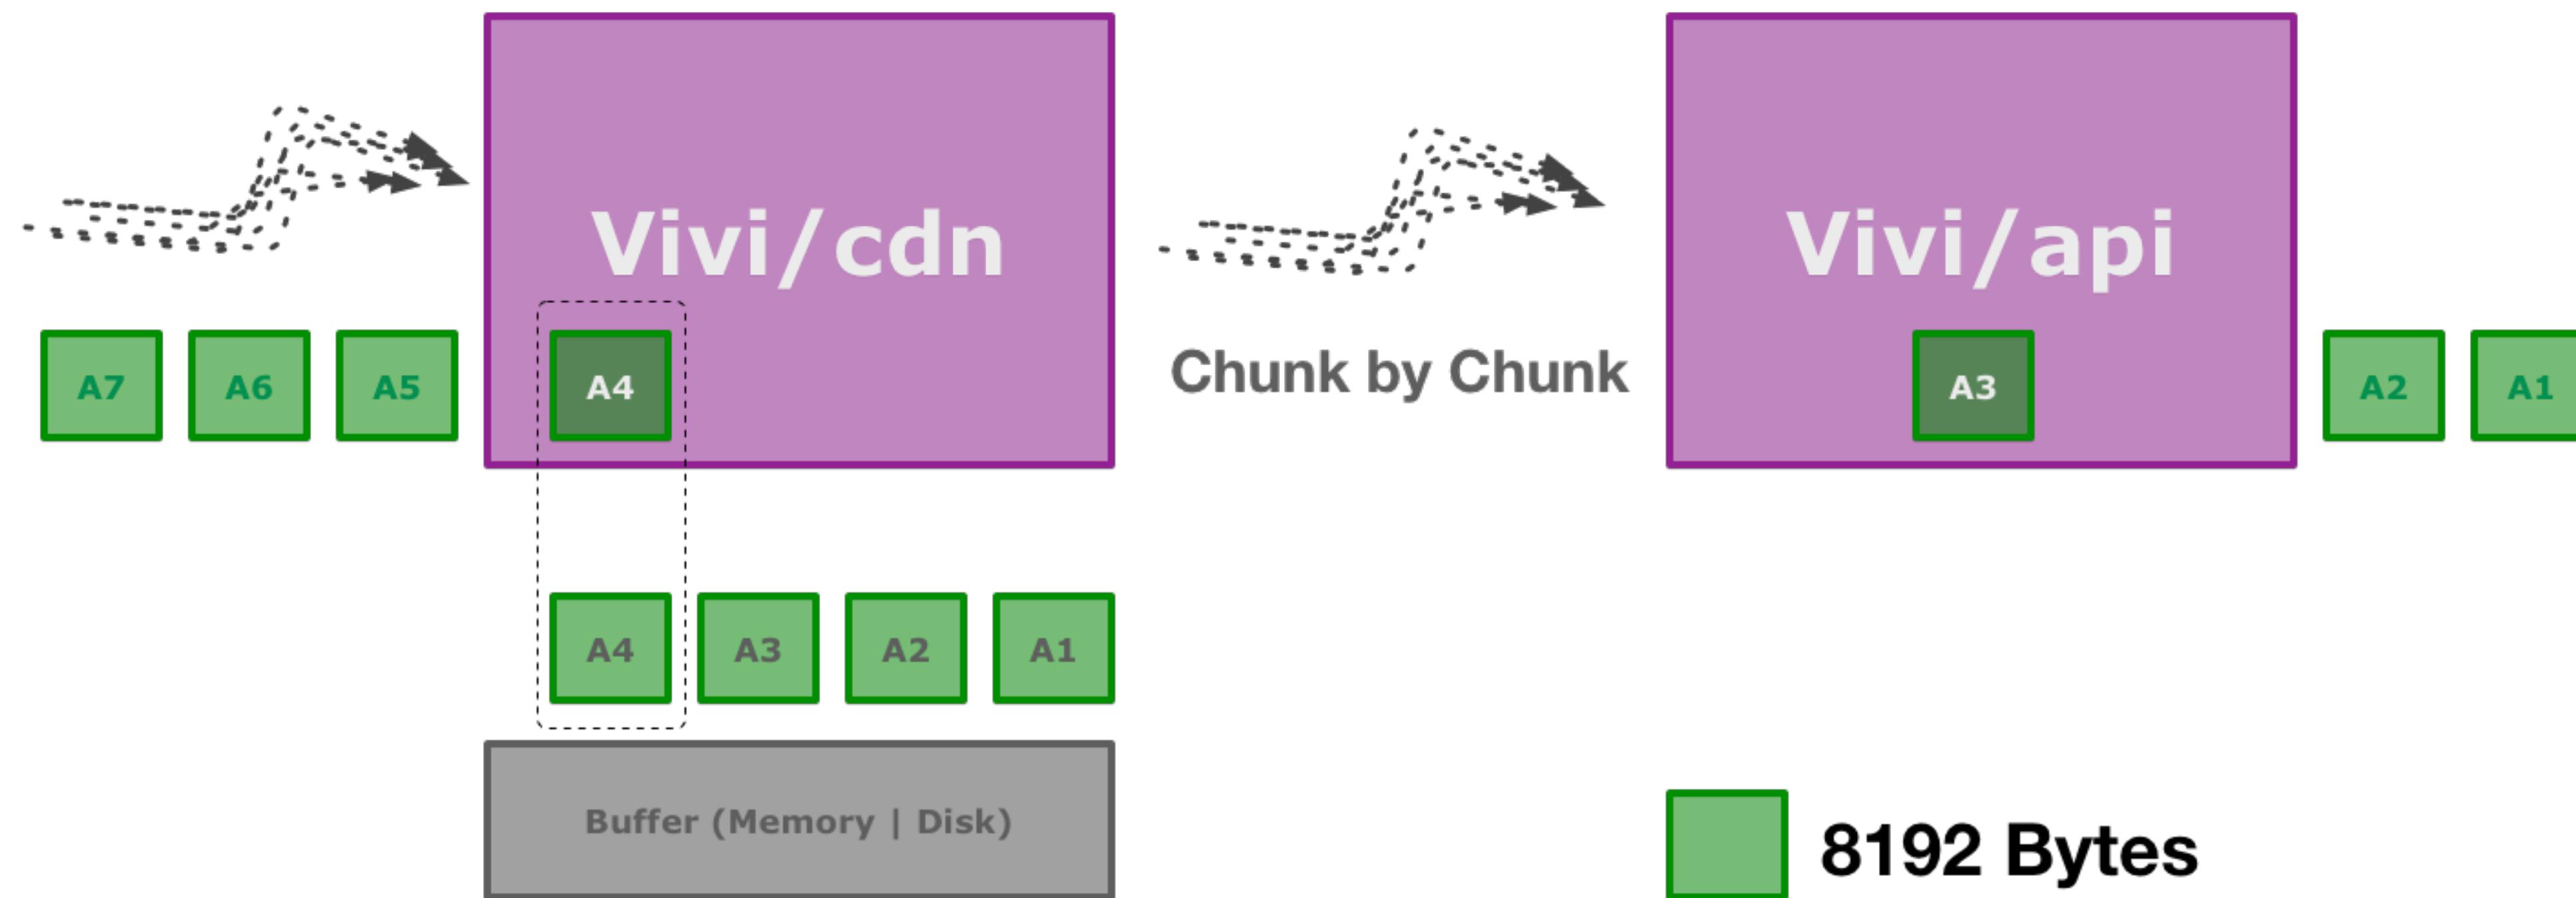


# UPYUN API



```
location /upload {  
    proxy_request_buffering off;  
    ...  
}
```

# Lua Streaming Upload



```
ngx.req.init_body()
```

```
ngx.req.append_body(chunk)
```

```
ngx.req.finish_body()
```

**Lua CDN**

**Lua WAF**

**Lua SSL**

**Lua API**

# Join our team



©2012-2014 Trybiane

Nginx `ngx_lua` agentzh `Lua`

[blog.cloudflare.com](http://blog.cloudflare.com) Openresty `LuaJIT` Github

Maxim Dounin Igor Sysoev chaoslawful

<https://groups.google.com/forum/#!forum/OpenResty>

Ansible Michael Pall Open source

# Thanks

...

Q & A